AIF-C01 Final Exam Cram Sheet

AWS Certified AI Practitioner - 30-45 Minute Pre-Exam Review

Task 12 of 15 | Read this 30-45 minutes before you walk into the exam. Do not start new material now - this is reinforcement only.

SECTION 1: EXAM QUICK FACTS

Fact	Detail
Total questions	65 (50 scored + 15 unscored)
Can you tell which are unscored?	No - answer every question as if it counts
Time allotted	90 minutes
Time per question	~83 seconds average (5,400 seconds / 65 questions)
Minimum passing score	700 out of 1,000 (scaled score, not a raw percentage)
Penalty for guessing	None - never leave a question blank
Question types	Multiple choice (1 correct, 3 distractors), multiple response ("Select TWO/THREE" - all correct choices required, no partial credit), ordering (arrange 3-5 steps), matching (pair 3-7 prompts to responses, all pairs required)
Domain 1 weight	20% (~10 scored questions) - Fundamentals of AI and ML
Domain 2 weight	24% (~12 scored questions) - Fundamentals of GenAI
Domain 3 weight	28% (~14 scored questions) - Applications of Foundation Models
Domain 4 weight	14% (~7 scored questions) - Guidelines for Responsible AI
Domain 5 weight	14% (~7 scored questions) - Security, Compliance, and Governance
Biggest domain	Domain 3 (28%) - foundation model applications is where the exam lives
Recertification	Every 3 years

> **Watch-out - services that sound right but are NOT in scope:** AWS Control Tower, Amazon Forecast, Amazon Fraud Detector, SageMaker Canvas, SageMaker Autopilot, and SageMaker Ground Truth are real AWS offerings but are **not** on the AIF-C01 in-scope services list. If an answer choice features one of these, treat it with extra suspicion - it is very likely a distractor. The exam also treats "Amazon Q" as a single generic AI assistant service rather than testing separate "Q Business" or "Q Developer" branding.

SECTION 2: HIGH-PRIORITY AWS SERVICES REFERENCE

Group 1 - Foundation Model and ML Platform Services

Service	One-Line Purpose	Top Exam Clue Phrase
Amazon Bedrock	Managed API access to multiple providers' foundation models without managing servers	"no ML expertise required to use a foundation model"
Amazon Bedrock Knowledge Bases	Connects an FM to your own data sources to ground responses (the RAG engine)	"connect the model to our company documents"
Amazon Bedrock Guardrails	Sets configurable safety limits on FM inputs and outputs (content filters, topic restrictions, PII redaction)	"block harmful or off-topic responses"
Amazon Bedrock Model Evaluation	AWS's managed service for running automated or human evaluation jobs against FMs	"compare model quality before choosing one"
Amazon Bedrock Prompt Management	Stores, versions, and manages prompts for production use	"track changes to a prompt over time"
Amazon Bedrock AgentCore	Managed runtime for deploying, securing, and scaling production AI agents	"deploy an agent securely at scale"
Amazon SageMaker AI	Full platform for building, training, and deploying custom ML models from scratch	"train our own custom model"
Amazon SageMaker JumpStart	Curated library of pre-built models and templates inside SageMaker AI	"pre-built model templates inside SageMaker"
Amazon SageMaker Clarify	Detects bias in data/models and explains feature importance	"detect bias" or "explain why the model made this prediction"
Amazon SageMaker Model Monitor	Tracks a deployed model's performance over time and alerts on drift	"is the model still performing well after deployment"
Amazon SageMaker Model Cards	Structured documentation of a model's intended use, limits, and training data	"document the model for stakeholders"
Amazon Augmented AI (A2I)	Routes low-confidence ML predictions to a human reviewer	"human review when the model is not confident"

Group 2 - Applied AI and Specialized Services

Service	One-Line Purpose	Top Exam Clue Phrase
Amazon Q	AI assistant for AWS questions, coding help, troubleshooting, and querying enterprise data	"AI assistant for employees/developers using AWS"
Kiro	AI-native IDE that uses agents to plan, write, and refine code	"AI-native development environment"
Strands Agents	Open-source framework for building agents (tools, memory, multi-step tasks)	"open-source agent-building framework"
Amazon Lex	Builds conversational chatbots and voice bots	"build a chatbot's dialogue flow"
Amazon Kendra	Intelligent enterprise search with natural-language Q&A across documents	"search across multiple internal data sources"
Amazon Comprehend	Analyzes text for sentiment, entities, and key phrases	"what does this text mean / how does the customer feel"
Amazon Textract	Extracts printed or handwritten text and form data from documents/images	"pull data out of a scanned form"
Amazon Transcribe	Converts spoken audio into text	"speech to text"
Amazon Polly	Converts written text into natural speech	"text to speech"
Amazon Translate	Neural machine translation between languages	"translate content into another language"
Amazon Rekognition	Analyzes images/video for objects, faces, text, and unsafe content	"what is in this picture or video"
Amazon Personalize	Builds real-time personalized recommendation systems	"customers also bought" / recommendation engine
Amazon Nova	Amazon's own family of foundation models, accessed via Bedrock	"Amazon-built foundation model"
AWS Transform	AI-powered legacy code modernization/migration assistance	"modernize legacy application code with AI"

Group 3 - Security and Governance Services

Service	One-Line Purpose	Top Exam Clue Phrase
AWS IAM	Controls who can access which AWS resources and what they can do	"control access permissions"

Service	One-Line Purpose	Top Exam Clue Phrase
AWS KMS	Creates and manages encryption keys for data at rest and in transit	"manage encryption keys"
AWS Secrets Manager	Securely stores and rotates application credentials (passwords, API keys)	"rotate a database password automatically"
Amazon Macie	ML-powered discovery of sensitive data (PII) stored in Amazon S3	"find PII in an S3 bucket"
Amazon Inspector	Continuously scans compute workloads for software vulnerabilities	"scan EC2/containers for vulnerabilities"
AWS PrivateLink	Private connectivity to AWS services without crossing the public internet	"keep traffic off the public internet"
AWS CloudTrail	Logs every API call and account action for auditing	"who did what and when"
Amazon CloudWatch	Monitors operational metrics, logs, and resource health	"system performance and health alerts"
AWS Config	Tracks AWS resource configuration state and compliance over time	"did this resource's configuration change"
AWS Artifact	Self-service portal to download AWS's own compliance reports/certifications	"download AWS's SOC 2 report"
AWS Audit Manager	Continuously collects evidence to support your own compliance audits	"build our own audit evidence automatically"
AWS Trusted Advisor	Automated, continuous recommendations across cost, security, performance	"real-time best-practice recommendations"
AWS Well-Architected Tool	Structured, question-driven review of a workload against best practices	"guided architecture review"

Vector Database Options for RAG (Domain 3 favorite)

Service	Why It Stores Vectors
Amazon OpenSearch Service	Purpose-built search/analytics engine; most common RAG vector store
Amazon Aurora	PostgreSQL-compatible with pgvector extension support
Amazon RDS for PostgreSQL	Same pgvector extension, within standard RDS
Amazon Neptune	Graph database that also supports vector search for connected data

SECTION 3: CRITICAL TERMINOLOGY LIGHTNING ROUND

One-sentence, exam-ready definitions. No elaboration.

AI and ML Fundamentals

- **Artificial Intelligence (AI):** The broad field of building systems that can perform tasks normally requiring human intelligence.
- **Machine Learning (ML):** A subset of AI where systems learn patterns from data instead of following explicit rules.
- **Deep Learning:** ML using multi-layered neural networks to learn complex patterns from large datasets.
- **Neural Network:** A model architecture loosely inspired by the brain, made of interconnected layers of nodes.
- **Algorithm:** The set of rules or mathematical steps a model follows to learn from data.
- **Model:** The trained artifact that takes input and produces output (predictions or generated content).
- **Training:** The process of teaching a model to find patterns in data.
- **Inferencing:** Using a trained model to make a prediction or generate output on new data.
- **Supervised Learning:** Learning from labeled data where the correct answer is already known.
- **Unsupervised Learning:** Finding patterns or groupings in unlabeled data with no known answer.
- **Reinforcement Learning:** Learning through trial and error guided by a reward signal.
- **Classification:** Assigning an input to one of several discrete categories.
- **Regression:** Predicting a continuous numerical value.
- **Clustering:** Grouping unlabeled data points into natural groups based on similarity.
- **Overfitting:** A model performs very well on training data but poorly on new, unseen data.
- **Underfitting:** A model performs poorly on both training data and new data.
- **Bias (statistical/ML):** Systematic error that skews model outputs in a consistent, unfair direction.
- **Accuracy / Precision / Recall / F1 Score:** Accuracy = overall correctness; Precision = of predicted positives, how many were correct; Recall = of actual positives, how many were caught; F1 = harmonic balance of precision and recall.

Generative AI and Foundation Models

- **Generative AI (GenAI):** AI that creates new content (text, image, audio, code) rather than just predicting a label.
- **Foundation Model (FM):** A large, pre-trained model adaptable to many different downstream tasks.
- **Large Language Model (LLM):** A foundation model specialized in understanding and generating human language.
- **Token:** A chunk of text (word or sub-word piece) that an LLM processes as a unit.
- **Chunking:** Breaking large documents into smaller pieces before embedding for retrieval.
- **Embedding:** A numeric vector representation of meaning used for similarity comparison.
- **Vector Database:** A database optimized to store and search embeddings by similarity.

- **Context Window:** The maximum amount of text (in tokens) a model can consider at one time.
- **Context Engineering:** Deliberately designing what information goes into the context window to improve output.
- **Multi-Modal Model:** A model that can process or generate more than one data type (e.g., text and images).
- **Hallucination:** A confident, fluent, but factually incorrect model output.
- **Nondeterminism:** The same prompt can produce a different output each time it is run.
- **Agentic AI:** AI systems that autonomously plan and take sequences of actions, calling tools to complete tasks.
- **AI Agent:** A single autonomous unit that can reason, use tools, and act toward a goal.
- **Model Context Protocol (MCP):** A standard way for agents to connect to external tools and data sources.
- **Token-Based Pricing:** Cost driven by the number of input and output tokens processed.
- **Provisioned Throughput:** Reserved, dedicated model capacity for predictable, high-volume workloads.

Foundation Model Applications and Techniques

- **Prompt Engineering:** The practice of designing inputs to guide a model toward a desired output.
- **Zero-Shot Prompting:** Asking the model to perform a task with no examples given.
- **Few-Shot Prompting:** Providing multiple examples in the prompt to guide the model's response format.
- **Chain-of-Thought Prompting:** Instructing the model to reason step by step before answering.
- **Retrieval Augmented Generation (RAG):** Retrieving relevant external data at inference time to ground a model's response, without changing model weights.
- **Fine-Tuning:** Further training an existing FM on a smaller, task-specific labeled dataset, which updates model weights.
- **In-Context Learning:** Improving output by including instructions/examples in the prompt itself, with no weight changes.
- **Model Distillation:** Training a smaller, faster model to mimic the outputs of a larger model.
- **RLHF (Reinforcement Learning from Human Feedback):** Using human preference ratings to further align a model's behavior.
- **ROUGE Score:** Automated metric for summarization quality (overlap with reference text).
- **BLEU Score:** Automated metric for translation quality (overlap with reference translation).
- **LLM-as-a-Judge:** Using a second, trusted LLM to evaluate the output quality of another model.
- **Grounding:** Anchoring a model's output in verified, retrieved, or factual source material.

Responsible AI

- **Fairness:** Equal and just treatment of all groups affected by an AI system.
- **Inclusivity:** Designing AI to work well across diverse populations and use cases.
- **Robustness:** Reliable model performance under varied or unexpected conditions.
- **Safety (AI):** Designing systems to avoid harmful outputs or actions.
- **Veracity:** The degree to which an AI system produces truthful, accurate outputs.

- **Transparency:** Being open about how a system works and what data/limitations it has.
- **Explainability:** The ability to explain why a model produced a specific decision or output.
- **Interpretability:** The degree to which a human can understand a model's internal logic.
- **Demographic Bias / Dataset Bias:** Skewed outputs caused by training data that misrepresents certain groups.
- **Model Card:** A structured document describing a model's intended use, limitations, and training data.

Security, Compliance, and Governance

- **Shared Responsibility Model:** AWS secures the cloud (infrastructure); the customer secures what they build in the cloud (data, access, configuration).
- **Data Lineage:** Tracking where data came from and how it has been transformed over time.
- **Data Governance:** The policies and processes that control how data is managed, accessed, and retained.
- **Prompt Injection:** An attack that manipulates a model's behavior through crafted user input.
- **Prompt Poisoning:** Malicious instructions injected into a model through tainted training or retrieved data.
- **Jailbreaking:** Attempting to bypass a model's built-in safety restrictions.
- **Data Leakage:** Sensitive information unintentionally appearing in a model's output.
- **Output Filtering / Validation:** Checking model outputs against defined rules or criteria before they reach the user.
- **Audit Trail:** A recorded, chronological history of actions taken within a system.
- **Generative AI Security Scoping Matrix:** An AWS framework that helps classify GenAI workloads by the level of responsibility a customer holds.

SECTION 4: THE MASTER COMPARISON TABLE

Service Pairs

Comparison	The One Thing to Remember
Amazon Bedrock vs. SageMaker AI	Bedrock = access pre-built FMs via API, no ML expertise needed. SageMaker AI = build, train, and deploy your own custom models.
Amazon Q vs. Kiro	Amazon Q = AI assistant for AWS questions, troubleshooting, and enterprise data, used within existing tools. Kiro = a full AI-native IDE built around agentic development workflows.
Amazon Lex vs. Amazon Kendra vs. Amazon Q	Lex = build chatbot dialogue flows. Kendra = intelligent search across enterprise documents. Amazon Q = broader AI assistant for AWS/coding/enterprise questions.

Comparison	The One Thing to Remember
Amazon Comprehend vs. Amazon Textract vs. Amazon Transcribe	Comprehend = analyze the meaning of text already in text form. Textract = extract text FROM images/scanned documents. Transcribe = convert spoken audio TO text.
Amazon Rekognition vs. Amazon Textract	Rekognition = analyze image/video content (what is in the picture). Textract = extract structured text and data from document images (what is written on the page).
SageMaker Clarify vs. Amazon A2I vs. SageMaker Model Monitor	Clarify = detect bias and explain model decisions. A2I = route low-confidence predictions to humans for review. Model Monitor = continuously track a deployed model's performance over time.
AWS Artifact vs. AWS Audit Manager	Artifact = download AWS's own pre-existing compliance reports. Audit Manager = continuously collect evidence for YOUR OWN compliance audits.
Amazon Macie vs. Amazon Inspector	Macie = scans S3 storage for sensitive data (PII). Inspector = scans compute workloads (EC2/containers) for software vulnerabilities.

Concept Pairs

Comparison	The One Thing to Remember
RAG vs. Fine-Tuning	RAG = retrieve external knowledge at inference time, model weights unchanged. Fine-Tuning = update model weights using domain-specific data, model changes.
Supervised vs. Unsupervised vs. Reinforcement Learning	Supervised = labeled data, known correct output. Unsupervised = no labels, find hidden patterns. Reinforcement = learn through trial, error, and reward signals.
Overfitting vs. Underfitting	Overfitting = excellent on training data, fails on new data. Underfitting = poor performance on both training and new data.
Hallucination vs. Bias	Hallucination = a confident, fluent, factually wrong output from the model itself. Bias = systematic unfairness rooted in skewed data or design choices.
Explainability vs. Interpretability vs. Transparency	Explainability = why did the model make THIS decision? Interpretability = can a human follow the model's internal logic in general? Transparency = is the system open about what it is, how it works, and its limitations?

Comparison	The One Thing to Remember
CloudTrail vs. CloudWatch vs. SageMaker Model Monitor	CloudTrail = who called which API and when (account activity log). CloudWatch = system health, metrics, and operational logs. Model Monitor = is my deployed ML model still performing well over time?
Encryption at Rest vs. In Transit	At rest = data stored on disk is encrypted. In transit = data moving across a network is encrypted (typically via TLS).

FM Customization Approaches at a Glance

Approach	Data Needed	Model Weights Change?	Best For
Prompt Engineering / In-Context Learning	None (just the prompt itself)	No	Quick, low-cost improvements with no training data available
RAG	A connected knowledge base or document source	No	Grounding answers in current, proprietary, or frequently changing data
Fine-Tuning	A labeled, task-specific dataset	Yes	Teaching the model a narrow skill, tone, or format the base model lacks
Continued Pre-Training	Large volume of unlabeled domain data	Yes	Deepening a model's general knowledge of an industry or domain over time

---

SECTION 5: RESPONSIBLE AI RAPID REVIEW

The 6 Official Features of Responsible AI (Task Statement 4.1)

Feature	One-Line Meaning
Bias	Systematic error that produces unfair or skewed results
Fairness	Equal and just treatment of all groups affected by the system
Inclusivity	Designing AI to work well for diverse populations and use cases
Robustness	Reliable performance under varied or unexpected conditions

Feature	One-Line Meaning
Safety	Avoiding harmful outputs or harmful actions
Veracity	Producing truthful, accurate outputs

Note: Transparency and Explainability are tested as their own dedicated topic under Task Statement 4.2, not bundled into the six features above - keep them mentally separate.

Dataset Characteristics That Support Responsible AI

- **Inclusivity:** representing diverse populations in the data.
- **Diversity:** covering varied scenarios and edge cases.
- **Curated/Vetted Sources:** data has been deliberately reviewed for quality and appropriateness.
- **Balanced Classes:** no single category dominates the dataset, preventing skewed learning.

Hallucination Mitigation Techniques (Task Statement 5.1)

- **RAG Grounding:** anchor responses in retrieved, verified source content.
- **Amazon Bedrock Guardrails:** enforce content filters, topic restrictions, and grounding checks.
- **Human Review (Amazon A2I):** route uncertain outputs to a human before they are acted on.
- **Output Validation:** check outputs against defined criteria before they reach the user.
- **Confidence Scoring:** surface a measure of how certain the model is about its own output.

SECTION 6: SECURITY RAPID REVIEW

Shared Responsibility for AI Services

	Amazon Bedrock	Amazon SageMaker AI
AWS manages	FM infrastructure, underlying compute, API availability, security "of" the cloud	Training/hosting infrastructure, managed notebook environment, platform security
Customer manages	Data sent to the model, IAM access control, logging, Guardrails configuration	Training data, model artifacts, endpoint access control, VPC/network configuration

> Per the AWS Shared Responsibility Model generally: **AWS is responsible for security OF the cloud** (host OS, virtualization layer, physical facilities). **The customer is responsible for security IN the cloud** (guest OS patching, application software, firewall/security group configuration, and IAM).

PII Protection - Two Tools, Two Different Moments

Tool	When It Acts	What It Does
Amazon Macie	On data already stored	Scans Amazon S3 to discover and classify sensitive data (PII, financial, health) at rest
Amazon Bedrock Guardrails (PII masking)	At inference time, in real time	Detects and masks/blocks PII in model inputs and outputs as a request is processed

---

SECTION 7: 30 EXAM CLUE PHRASE PAIRS

Domain 1 - Fundamentals of AI and ML

- If you see "convert spoken audio into text"… think **Amazon Transcribe**.
- If you see "convert written text into natural speech"… think **Amazon Polly**.
- If you see "translate content between languages"… think **Amazon Translate**.
- If you see "detect sentiment, entities, or key phrases in text"… think **Amazon Comprehend**.
- If you see "predicting a continuous numeric value like price or temperature"… think **regression**.
- If you see "group unlabeled data into natural categories with no predefined labels"… think **clustering / unsupervised learning**.

Domain 2 - Fundamentals of GenAI

- If you see "how much text a model can consider at once"… think **context window**.
- If you see "the same prompt produces a different answer each time it runs"… think **nondeterminism**.
- If you see "breaking a long document into smaller pieces before embedding"… think **chunking**.
- If you see "a numeric representation of meaning used for similarity search"… think **embedding / vector**.
- If you see "cost scales with the amount of input and output text processed"… think **token-based pricing**.
- If you see "an agent calling an external tool or data source through a standard interface"… think **Model Context Protocol (MCP)**.

Domain 3 - Applications of Foundation Models

- If you see "connect a foundation model to internal company documents without retraining it"… think **RAG, Amazon Bedrock Knowledge Bases**.
- If you see "controls how creative versus predictable the model's output is"… think **temperature (inference parameter)**.
- If you see "a few examples included directly inside the prompt"… think **few-shot prompting**.
- If you see "ask the model to reason step by step before giving a final answer"… think **chain-of-thought prompting**.
- If you see "a smaller, faster model trained to mimic a larger model's outputs"… think **model distillation**.
- If you see "compare machine translation output against a human reference translation"… think **BLEU score**.

Domain 4 - Guidelines for Responsible AI

- If you see "model performs perfectly on training data but fails on new data"… think **overfitting**.
- If you see "detect bias in a training dataset or a deployed model"… think **Amazon SageMaker Clarify**.
- If you see "document a model's intended use, limitations, and training data for stakeholders"… think **Amazon SageMaker Model Cards**.
- If you see "route a low-confidence prediction to a human reviewer"… think **Amazon Augmented AI (A2I)**.
- If you see "explain why the model produced this specific decision"… think **explainability**.
- If you see "training data does not represent all demographic groups equally"… think **dataset bias / unbalanced dataset**.

Domain 5 - Security, Compliance, and Governance

- If you see "who made this API call" or "a log of account activity"… think **AWS CloudTrail**.
- If you see "find personally identifiable information stored in an S3 bucket"… think **Amazon Macie**.
- If you see "malicious instructions hidden inside user input or retrieved data"… think **prompt injection**.
- If you see "download AWS's own SOC 2 report or compliance certification"… think **AWS Artifact**.
- If you see "continuously collect evidence to support our own internal compliance audit"… think **AWS Audit Manager**.
- If you see "private connectivity to an AWS service without traversing the public internet"… think **AWS PrivateLink**.

SECTION 8: EXAM-DAY TACTICS

1. **Flag and return.** If a question is taking too long, flag it, make your best guess, and move on. Never let one question burn more than ~90 seconds.
2. **Read the qualifier word carefully.** "Most appropriate," "most cost-effective," and "most secure" are three different questions that can have three different correct answers for the exact same scenario.
3. **For "Select TWO/THREE" questions, evaluate all options independently first.** Decide whether each option is true or false on its own merits before comparing them against each other.
4. **Use process of elimination aggressively.** Once you know the material, wrong answers are usually wrong for an obvious, specific reason - find that reason rather than searching for the "best" answer from scratch.
5. **For long scenario questions, read the final sentence (the actual question) first.** Then read the scenario already knowing what you are being asked to find.
6. **"No ML expertise required" is a strong signal for managed AI services** like Amazon Bedrock, Amazon Q, or Amazon Rekognition - not SageMaker AI.
7. **"Custom model" or "train your own" is a strong signal for Amazon SageMaker AI**, not Bedrock.

8. **Trust your preparation.** If you studied the material, your first instinct on a question is usually correct - resist the urge to second-guess yourself into a wrong answer.
9. **Never leave a question blank.** There is no penalty for an incorrect answer, but a blank answer is automatically scored as wrong.
10. **Treat every question as scored.** The 15 unscored pilot questions are not flagged in any way - answer everything with full effort.

SECTION 9: THINGS TO MEMORIZE vs. THINGS TO UNDERSTAND

Worth Memorizing (exact facts)	Worth Understanding (how/why it works)
65 total questions, 50 scored / 15 unscored	How RAG retrieves and injects context into a prompt
90-minute time limit	Why the Shared Responsibility Model shifts as you move from EC2 to Bedrock
700/1000 passing score (scaled)	When to choose fine-tuning over RAG (and vice versa)
Domain weights: 20/24/28/14/14	How hallucinations happen and why grounding reduces them
The 6 features of responsible AI (bias, fairness, inclusivity, robustness, safety, veracity)	What makes a training dataset biased or unbalanced
Amazon Bedrock = pre-built FMs via API	Why temperature affects creativity vs. predictability of output
Amazon SageMaker AI = build custom models	The tradeoff between model transparency and model performance
Macie = S3 sensitive-data scanning	How prompt injection differs from prompt poisoning differs from jailbreaking
Inspector = compute vulnerability scanning	Why token-based pricing rewards shorter, well-engineered prompts
CloudTrail = API activity logging	How few-shot prompting differs in effect from chain-of-thought prompting
Artifact = download AWS's compliance reports	Why overfitting happens and what it looks like in evaluation metrics
Audit Manager = build your own audit evidence	How A2I, Clarify, and Model Monitor form a responsible-AI lifecycle together
ROUGE = summarization metric, BLEU = translation metric	What context engineering actually controls in an FM application
RAG = no weight change; Fine-tuning = weight change	Why agentic AI needs memory, tool usage, and orchestration to complete multi-step tasks
4 question types: multiple choice, multiple response, ordering, matching	How the Generative AI Security Scoping Matrix helps classify a workload's responsibility level

---

End of Task 12 - Final Exam Cram Sheet Use this document in the final 30-45 minutes before your scheduled AIF-C01 exam session. Good luck.