AIF-C01 Master Project Blueprint

AWS AI Cloud Practitioner Certification Prep Program

Project Reference Document - Task 1 of 15 Version 1.0 | Based on AIF-C01 Exam Guide Version 1.1 (April 30, 2026)

SECTION 1: EXAM OVERVIEW

Official Exam Identity

Field	Detail
Official Exam Name	AWS Certified AI Practitioner
Exam Code	AIF-C01
Certification Tier	Foundational
Total Questions	65 (50 scored + 15 unscored)
Time Allotted	90 minutes
Minimum Passing Score	700 out of 1,000 (scaled score)
Delivery Method	Pearson VUE - testing center or online proctored
Recertification Period	3 years
Current Exam Guide Version	1.1 (published April 30, 2026)

Question Types

Question Type	Description	Key Rule
Multiple Choice	One correct answer, three distractors (incorrect options)	Eliminate clearly wrong answers before choosing
Multiple Response	Two or more correct answers out of five or more options; all correct choices must be selected	No partial credit - all correct options must be chosen
Ordering	A list of 3-5 steps that must be arranged in the correct sequence	Think about logical process flow
Matching	Match 3-7 prompts to their corresponding responses; all pairs must be correct	No partial credit - every pair must match

Scoring note: Unanswered questions are scored as incorrect. There is no penalty for guessing, so always select an answer even when uncertain.

Scored vs. Unscored Questions

The exam contains 65 total questions. Fifty of these questions affect your score. The remaining 15 are unscored pilot questions that AWS uses to evaluate potential future exam questions. You will not be told which questions are unscored, so treat every question as though it counts.

What This Certification Validates

The AWS Certified AI Practitioner (AIF-C01) validates that a candidate can:

- Describe AI, ML, and generative AI (GenAI) concepts, methods, and strategies in general and within the AWS ecosystem.
- Identify the appropriate use of AI/ML and GenAI technologies to solve real business problems.
- Determine which types of AI/ML technologies apply to specific business scenarios.
- Use AI, ML, and GenAI technologies responsibly, with awareness of ethical considerations, security requirements, and governance frameworks.

The target candidate has up to six months of exposure to AI/ML technologies on AWS. Importantly, this exam is for people who **use** AI/ML solutions - not those who **build or code** them from scratch. Candidates are not expected to write code, tune models mathematically, or deploy infrastructure.

What Is Out of Scope for This Exam

The following tasks are NOT tested on AIF-C01:

- Developing or coding AI/ML models or algorithms
- Implementing data engineering or feature engineering techniques
- Performing hyperparameter tuning or model optimization
- Building and deploying AI/ML pipelines or infrastructure from scratch
- Conducting mathematical or statistical analysis of AI/ML models
- Implementing security or compliance protocols from a developer perspective
- Creating governance frameworks and policies from scratch

Why This Certification Matters

AI and machine learning capabilities are reshaping every industry. Organizations that understand how to apply these technologies gain competitive advantages in automation, customer experience, decision-making speed, and operational efficiency. The AIF-C01 certification signals that a professional:

- Can have informed conversations with technical teams about AI solutions
- Understands how to evaluate and select AI tools for business problems
- Is aware of the risks, ethical considerations, and compliance requirements that come with AI adoption
- Knows the AWS AI/ML ecosystem well enough to participate in procurement, planning, and governance decisions

For learners in business, operations, or junior technical roles, this certification demonstrates foundational AI literacy that is increasingly expected across industries - from healthcare and finance to retail and manufacturing. As AI tools become standard workplace tools, professionals who understand them at a conceptual and practical level hold a significant career advantage.

SECTION 2: OFFICIAL DOMAIN MAP

Domain Summary Table

Domain #	Domain Name	Weight %	Scored Questions (~)	Plain-English Focus	Top In-Scope AWS Services	Learner Difficulty
1	Fundamentals of AI and ML	20%	~10	What AI and ML are, how they work at a basic level, what types of problems they solve, and how the development lifecycle works. Builds the vocabulary used in all other domains.	Amazon SageMaker AI, Amazon Comprehend, Amazon Rekognition, Amazon Transcribe, Amazon Polly	Moderate - primarily vocabulary-heavy
2	Fundamentals of GenAI	24%	~12	What generative AI is, how foundation models and large language models work, the advantages and limitations of GenAI for business, and the AWS tools available for building GenAI applications.	Amazon Bedrock, Amazon SageMaker AI, Amazon Q, Strands Agents, Amazon Bedrock AgentCore	Moderate to High - fast-moving topic area with new terminology

Domain #	Domain Name	Weight %	Scored Questions (~)	Plain-English Focus	Top In-Scope AWS Services	Learner Difficulty
3	Applications of Foundation Models	28%	~14	How to select, configure, and customize foundation models; how to write effective prompts; how to enhance models with external data (RAG); and how to measure whether a model is actually working well.	Amazon Bedrock, Amazon Bedrock Knowledge Bases, Amazon OpenSearch Service, Amazon Bedrock Guardrails, Amazon Bedrock Model Evaluation	High - most heavily weighted; requires applying concepts, not just recalling them

Domain #	Domain Name	Weight %	Scored Questions (~)	Plain-English Focus	Top In-Scope AWS Services	Learner Difficulty
4	Guidelines for Responsible AI	14%	~7	The ethical principles behind AI development - fairness, bias, transparency, and explainability - and the AWS tools that help organizations build AI systems that are safe and trustworthy.	Amazon SageMaker Clarify, Amazon SageMaker Model Cards, Amazon A2I, Amazon Bedrock Guardrails, SageMaker Model Monitor	Moderate - conceptual but closely tied to specific AWS tools
5	Security, Compliance, and Governance for AI Solutions	14%	~7	How to protect AI systems from threats, how to meet regulatory requirements, and how to establish governance structures that keep AI use safe, auditable, and compliant.	AWS IAM, AWS CloudTrail, AWS Artifact, AWS Audit Manager, Amazon Macie	Moderate - builds on general AWS security knowledge

Domain 1: Fundamentals of AI and ML (20%)

Domain 1 is the starting point for the entire certification. Think of it as building the dictionary you will use for everything else you study. Before you can talk about how to use AI tools wisely, you need to understand what AI actually is, what types of learning exist, and what kinds of problems AI can and cannot solve.

This domain covers three broad areas. First, it asks you to define key vocabulary: terms like algorithm, model, supervised learning, neural network, and inferencing. Second, it asks you to recognize which AI techniques fit which business scenarios - for example, knowing that classification is useful for detecting fraudulent transactions, while regression works better for predicting sales numbers. Third, it introduces the AI/ML development lifecycle: the journey from collecting raw data all the way through deploying and monitoring a model in production.

A learner who has mastered Domain 1 can explain AI concepts clearly to a non-technical colleague, identify when a business problem is a good candidate for an AI solution (and when it is not), describe the stages an AI/ML project goes through from start to finish, and name the key AWS managed AI services along with what they do.

Domain 2: Fundamentals of GenAI (24%)

Domain 2 focuses on generative AI - the technology behind chatbots, AI writing assistants, image generators, and code completion tools. This is one of the most rapidly evolving areas in technology, and this domain gives learners a foundational understanding of how it works, what makes it powerful, and where it falls short.

Generative AI works differently from traditional AI. Instead of predicting a single answer (like "is this email spam?"), it generates new content - text, images, code, audio - based on patterns learned from massive amounts of training data. This domain explains the building blocks of that process: tokens (the chunks of text that models process), embeddings (how meaning is represented mathematically), and context windows (how much information a model can consider at once). It also introduces agentic AI - systems where AI can take sequences of actions autonomously, calling tools and making decisions to complete complex tasks.

On the business side, Domain 2 asks learners to weigh GenAI's advantages (speed, adaptability, the ability to handle many types of tasks) against its disadvantages (hallucinations, unpredictability, cost). It also maps out the AWS tools - primarily Amazon Bedrock - that allow organizations to build GenAI applications without needing to train massive models from scratch.

Domain 3: Applications of Foundation Models (28%)

Domain 3 is the most heavily weighted domain on the exam and the most practically focused. It is about making foundation models work well for real business purposes. A foundation model (FM) is a large, pre-trained AI model that can be adapted for many different tasks - but using one effectively requires deliberate design decisions.

This domain covers four main areas: how to choose the right FM for a given use case (based on factors like cost, speed, language support, and output quality); how to write effective prompts (the instructions you give an AI model to guide its responses); how to customize or teach a model new things through fine-tuning; and how to measure whether a model is actually doing its job well. A major topic in this domain is Retrieval Augmented Generation (RAG) - a technique that lets an AI model search a company's own documents or knowledge base before responding, producing more accurate and grounded answers.

Learners who master Domain 3 can describe the tradeoffs between different ways of customizing an AI model, write more effective prompts using proven techniques, explain what RAG is and why it matters for enterprise use cases, and identify the specific metrics used to assess whether a model is meeting business goals.

Domain 4: Guidelines for Responsible AI (14%)

Domain 4 addresses the ethical and human side of AI development. As AI systems become more capable and more widely deployed, the potential for unintended harm grows alongside their benefits. This domain ensures that learners understand not just what AI can do, but what guardrails are needed to ensure it acts safely and fairly.

Key topics include fairness (ensuring AI does not systematically disadvantage certain groups of people), bias (recognizing when training data or model behavior produces skewed or unfair results), and transparency and explainability (understanding and being able to explain why an AI made a particular decision). This domain also covers legal risks - such as copyright concerns with AI-generated content - and environmental considerations around large-scale AI training and deployment.

AWS provides specific tools to help with these challenges, including Amazon SageMaker Clarify (for detecting bias in models and datasets), Amazon SageMaker Model Cards (for documenting and communicating model behavior), and Amazon Bedrock Guardrails (for setting safety limits on what a model can say or do). Learners in Domain 4 need to know what these tools do and when to use them - without needing to implement them from scratch.

Domain 5: Security, Compliance, and Governance for AI Solutions (14%)

Domain 5 brings together security awareness and regulatory knowledge in the context of AI systems. Organizations deploying AI face security risks that go beyond traditional software, including prompt injection attacks (attempts to manipulate an AI through cleverly crafted input), data leakage (sensitive information accidentally appearing in model outputs), and the challenge of keeping AI behavior auditable and aligned with regulations.

This domain covers how familiar AWS security services - like IAM, AWS KMS, Amazon Macie, and AWS PrivateLink - apply specifically to AI workloads. It also covers governance tools like AWS CloudTrail, AWS Artifact, and AWS Audit Manager that help organizations demonstrate compliance to auditors and regulators. The AWS Shared Responsibility Model is central here: AWS secures the underlying infrastructure, while customers remain responsible for what they build and configure on top of it.

Learners who succeed in Domain 5 understand the types of threats unique to AI systems, know which AWS services provide security and compliance support for AI workloads, and can describe the governance processes an organization should follow when deploying AI solutions responsibly.

SECTION 3: COMPLETE CONTENT INVENTORY

Program Deliverables Table

Task #	Deliverable Name	Type	Primary Audience	Phase	Domain(s)	Depends On
1	Master Project Blueprint	Reference Doc	Instructor / Admin	Foundation	All	None
2	Learner Orientation and Program Overview Guide	Support Material	Learner	Foundation	All	Task 1
3	Domain 1 Study Module: Fundamentals of AI and ML	Study Module	Learner	Domain Modules	Domain 1	Task 2
4	Domain 2 Study Module: Fundamentals of GenAI	Study Module	Learner	Domain Modules	Domain 2	Task 3
5	Domain 3 Study Module: Applications of Foundation Models	Study Module	Learner	Domain Modules	Domain 3	Task 4
6	Domain 4 Study Module: Guidelines for Responsible AI	Study Module	Learner	Domain Modules	Domain 4	Task 5
7	Domain 5 Study Module: Security, Compliance, and Governance for AI Solutions	Study Module	Learner	Domain Modules	Domain 5	Task 6

Task #	Deliverable Name	Type	Primary Audience	Phase	Domain(s)	Depends On
8	AWS Services Quick Reference Guide	Reference Doc	Learner	Reference Materials	All	Task 7
9	Master Glossary	Reference Doc	Learner	Reference Materials	All	Tasks 3-7
10	Visual Cheat Sheet and One-Page Summary	Reference Doc	Learner	Reference Materials	All	Tasks 3-7
11	Master Flashcard Deck	Assessment	Learner	Assessment	All	Tasks 3-7
12	Domain Practice Question Banks (All Five Domains)	Assessment	Learner	Assessment	All	Tasks 3-7
13	Full-Length Practice Exam (65 Questions)	Assessment	Learner	Assessment	All	Task 12
14	Practice Exam Answer Key and Detailed Rationale Guide	Assessment	Learner	Assessment	All	Task 13
15	Exam Day Strategy Guide and Final Readiness Checklist	Support Material	Learner	Final Output	All	Tasks 1-14

Deliverable Descriptions

Task 1 - Master Project Blueprint: The foundational reference document for the entire program. Contains the domain map, service registry, term list, task statement analysis, and recommended study sequence. This document is the single source of truth for all subsequent tasks and is used by the instructor to maintain consistency across all deliverables.

Task 2 - Learner Orientation Guide: An accessible introduction that orients learners to the exam, sets expectations for difficulty and time investment, explains the study program structure, and provides a recommended weekly study schedule template. Written in plain language with minimal assumptions about prior technical knowledge.

Tasks 3-7 - Domain Study Modules: One comprehensive study module per exam domain. Each module covers all task statements for that domain using plain-language explanations, real-world analogies, examples, AWS service descriptions, exam watch-outs, and embedded knowledge check questions. Each module is self-contained and can stand alone as a reference after the program ends.

Task 8 - AWS Services Quick Reference Guide: A focused reference document covering every in-scope AWS service organized by category and exam priority. Includes concise descriptions, primary use cases, relevant exam domains, and clear explanations of how commonly confused services differ. Designed for review and memorization support, not deep study.

Task 9 - Master Glossary: An alphabetical dictionary of every key term used across all five domains. Each entry provides a plain-language definition, the domain it belongs to, and a real-world analogy where applicable. Designed for use as a standalone reference throughout the entire study program and for final review before the exam.

Task 10 - Visual Cheat Sheet: A compact, highly structured one-to-two page summary of the most critical facts, service names, term definitions, and relationships across all five domains. Designed for quick review in the days immediately before the exam. Formatted for printing.

Task 11 - Master Flashcard Deck: A comprehensive set of flashcards covering all high-priority and medium-priority terms, AWS services, and exam concepts. Each card follows a consistent format with the term or question on one side and the definition, context, and exam tip on the other. Organized by domain and by priority level.

Task 12 - Domain Practice Question Banks: A set of original practice questions organized by domain, with a distribution proportional to the exam weighting (approximately 10 for Domain 1, 12 for Domain 2, 14 for Domain 3, 7 for Domain 4, and 7 for Domain 5). Questions span all four exam question types. Each question includes the correct answer and a brief rationale.

Task 13 - Full-Length Practice Exam: A complete 65-question simulated exam that mirrors the AIF-C01 format, timing, and domain distribution. Includes all four question types. Presented as a standalone timed exam experience, separate from the answer key.

Task 14 - Practice Exam Answer Key and Rationale Guide: Detailed explanations for every question in the full-length practice exam (Task 13). Each explanation identifies the correct answer, explains why it is correct, explains why each incorrect option is wrong, and references the relevant domain and task statement. The most valuable study document for learners who have already taken the practice exam.

Task 15 - Exam Day Strategy Guide: A practical guide covering test-taking strategies specific to all four AIF-C01 question types, time management techniques for 90 minutes across 65 questions, common traps and distractor patterns to avoid, and a day-of readiness checklist covering logistics, equipment (for online proctored delivery), and mental preparation.

SECTION 4: IN-SCOPE AWS SERVICES REGISTRY

Machine Learning Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Bedrock	Machine Learning	Domains 2, 3, 4, 5	High	The primary AWS platform for accessing, customizing, and deploying foundation models from multiple providers through a single managed API - without managing servers.	Amazon SageMaker AI: Bedrock is for using pre-built FMs; SageMaker AI is for building and training custom models from scratch.
Amazon Bedrock AgentCore	Machine Learning	Domains 2, 3, 5	High	A managed service for building, deploying, securing, and scaling production-grade AI agents that can take actions, call tools, and orchestrate multi-step workflows.	Strands Agents: AgentCore is the managed production runtime for agents; Strands Agents is the open-source development framework used to build agents.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon SageMaker AI	Machine Learning	Domains 1, 3, 4	High	The comprehensive AWS platform for building, training, and deploying custom machine learning models at any scale, from experimentation through production.	Amazon Bedrock: SageMaker AI is for building and training your own models; Bedrock is for using existing foundation models via API.
Amazon SageMaker JumpStart	Machine Learning	Domains 1, 2, 3	High	A feature within SageMaker that provides a curated library of pre-built ML models and solution templates that can be deployed with minimal setup, including select foundation models.	Amazon Bedrock: JumpStart offers FMs accessed through the SageMaker interface; Bedrock is the dedicated managed FM platform with broader model selection.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Augmented AI (Amazon A2I)	Machine Learning	Domain 4	High	A service that adds human review workflows into ML prediction pipelines, used when a model's confidence is low and a human expert should verify the output before acting on it.	SageMaker Model Monitor: A2I routes low-confidence predictions to human reviewers; Model Monitor automatically tracks and alerts on model performance degradation.
Amazon Comprehend	Machine Learning	Domain 1	High	A natural language processing (NLP) service that analyzes text to detect sentiment (positive, negative, neutral), identify entities (names, dates, places), classify topics, and extract key phrases.	Amazon Textract: Comprehend analyzes the meaning and content of text that is already in text form; Textract extracts printed or handwritten text from images and scanned documents.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Kendra	Machine Learning	Domain 3	Medium	An intelligent enterprise search service powered by ML that allows employees to search across multiple data sources using natural language and receive precise, relevant answers.	Amazon OpenSearch Service: Kendra is purpose-built for intelligent enterprise document search with out-of-the-box NLP; OpenSearch is a flexible general-purpose search and analytics engine often used to store vectors for RAG.
Amazon Lex	Machine Learning	Domain 1	High	A service for building conversational interfaces - chatbots and voice bots - that can understand natural language, powered by the same AI technology that drives Amazon Alexa.	Amazon Polly: Lex understands and responds to conversation (speech-to-intent); Polly converts written text into spoken audio (text-to-speech).

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Nova	Machine Learning	Domains 2, 3	High	Amazon's own family of foundation models, available through Amazon Bedrock, including text, image, and multi-modal models designed for high performance and cost-efficiency.	Amazon Bedrock (platform): Nova refers to specific Amazon-built foundation models; Bedrock is the platform through which Nova models and third-party models are all accessed.
Amazon Personalize	Machine Learning	Domain 1	Medium	A service that uses ML to build real-time personalized recommendation systems - similar to the engines that power "customers also bought" and "recommended for you" features.	Amazon Comprehend: Personalize generates recommendations based on user behavior data; Comprehend analyzes the content and sentiment of text.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Polly	Machine Learning	Domain 1	High	A text-to-speech service that converts written text into natural-sounding spoken audio, supporting dozens of languages and voice styles including neural voices.	Amazon Transcribe: Polly converts text into speech (text-to-speech); Transcribe converts spoken audio into text (speech-to-text). These are opposites.
Amazon Rekognition	Machine Learning	Domain 1	High	A computer vision service that analyzes images and video to detect and label objects, faces, text, scenes, activities, inappropriate content, and custom categories.	Amazon Textract: Rekognition analyzes visual content in images and video (what is in the picture); Textract extracts structured text and data from document images (what is written on the page).

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Textract	Machine Learning	Domain 1	Medium	A service that automatically extracts printed text, handwriting, and structured data (like form fields and table values) from scanned documents, images, and PDFs.	Amazon Comprehend: Textract extracts text from document images into a usable format; Comprehend then analyzes the meaning and content of that extracted text.
Amazon Transcribe	Machine Learning	Domain 1	High	An automatic speech recognition (ASR) service that converts spoken audio - from calls, recordings, or live streams - into accurate written text transcripts.	Amazon Polly: Transcribe converts speech to text; Polly converts text to speech. They work in opposite directions.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Translate	Machine Learning	Domain 1	High	A neural machine translation service that provides fast, high-quality language translation between dozens of supported languages, used for content localization and multilingual applications.	Amazon Comprehend: Translate converts text from one language to another; Comprehend analyzes text for meaning, entities, and sentiment within a single language.
AWS Transform	Machine Learning	Domain 1	Medium	An AI-powered service that helps organizations modernize and migrate legacy application code, accelerating digital transformation efforts through automated code analysis and conversion.	AWS Migration Hub: Transform uses AI to modernize and convert application code; Migration Hub is a tracking and orchestration service for broader cloud migration projects.

Analytics Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
AWS Data Exchange	Analytics	Domain 1	Low	A marketplace where organizations can discover, subscribe to, and use third-party datasets directly within AWS for analytics and ML model training.	AWS Glue: Data Exchange is for acquiring licensed external data; Glue is for preparing, transforming, and loading data once you have it.
Amazon EMR	Analytics	Domain 1	Low	A managed big data platform that runs open-source frameworks (Apache Spark, Hadoop) to process very large datasets in parallel, used in data preparation pipelines for ML.	AWS Glue: EMR is a full cluster-based big data processing platform requiring more configuration; Glue is a simpler, serverless ETL service better suited to standard data integration tasks.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
AWS Glue	Analytics	Domain 1	Medium	A serverless data integration service that discovers, prepares, and transforms data from multiple sources, commonly used in the data preparation stage of ML pipelines.	Amazon EMR: Glue is simpler and serverless (no cluster to manage); EMR is a more powerful but more complex cluster-based processing platform.
AWS Glue DataBrew	Analytics	Domain 1	Low	A visual, no-code data preparation tool that allows business analysts and data scientists to clean and normalize data without writing code, as an entry point to ML data preparation.	AWS Glue: DataBrew is the visual, point-and-click version of data preparation; Glue is the code-based serverless ETL service for engineers.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
AWS Lake Formation	Analytics	Domain 1	Low	A service that helps organizations quickly build, secure, and manage a central data lake in Amazon S3, with fine-grained access controls, data cataloging, and governance features.	Amazon S3: S3 stores the raw data; Lake Formation adds a governance, cataloging, and access control layer on top of that storage.
Amazon OpenSearch Service	Analytics	Domain 3	High	A managed search and analytics engine that can store and query vector embeddings, making it a common choice for building the vector database component of RAG-based applications.	Amazon Kendra: OpenSearch is a flexible search and analytics engine used for vector storage, log analysis, and custom search; Kendra is a higher-level managed intelligent search service designed specifically for enterprise Q-and-A over documents.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Quick	Analytics	Domains 1, 2	Medium	A business intelligence and data visualization service that enables organizations to create interactive dashboards and reports, with AI-powered natural language query capabilities for non-technical users.	Amazon Redshift: Quick (formerly aligned with QuickSight) is the front-end BI and visualization layer; Redshift is the cloud data warehouse where the underlying analytical data is stored and queried.
Amazon Redshift	Analytics	Domain 1	Low	A fully managed cloud data warehouse optimized for analytical queries across large volumes of structured data, commonly used as a data source for ML training datasets and business reporting.	Amazon RDS: Redshift is an analytical data warehouse designed for complex queries over large historical datasets; RDS is an operational relational database designed for transactional workloads.

Security, Identity, and Compliance Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
AWS Identity and Access Management (IAM)	Security	Domains 4, 5	High	The foundational AWS service for controlling who (users, roles, services) can access which AWS resources and what actions they can perform, using policies and permissions.	AWS Secrets Manager: IAM controls access to AWS services and resources; Secrets Manager securely stores and retrieves application credentials like database passwords and API keys.
AWS Artifact	Security	Domain 5	High	A self-service portal where AWS customers can access and download AWS compliance reports, security certifications, and audit documentation on demand, at no cost.	AWS Audit Manager: Artifact is where customers download AWS's own compliance reports and certifications; Audit Manager helps customers build and maintain their own internal compliance evidence.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
AWS Audit Manager	Security	Domain 5	High	A service that continuously and automatically collects evidence of AWS resource configurations and activity, helping organizations demonstrate compliance with frameworks like SOC 2, GDPR, and HIPAA.	AWS Artifact: Audit Manager is used by customers to conduct and document their own ongoing compliance audits; Artifact provides pre-existing AWS audit reports for customers to download.
Amazon Inspector	Security	Domain 5	Medium	An automated vulnerability management service that continuously scans AWS compute workloads (EC2 instances, containers) for known software vulnerabilities and unintended network exposure.	Amazon Macie: Inspector scans compute workloads for security vulnerabilities; Macie scans Amazon S3 storage for sensitive data that should not be exposed.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
AWS Key Management Service (AWS KMS)	Security	Domain 5	Medium	A managed service for creating, storing, and controlling the encryption keys used to protect data at rest and in transit across AWS services and customer applications.	AWS Secrets Manager: KMS manages encryption keys (the keys used to lock and unlock data); Secrets Manager manages application secrets like passwords and API tokens that applications need to authenticate.
Amazon Macie	Security	Domains 4, 5	High	An ML-powered data security service that automatically discovers and classifies sensitive data (such as PII, financial records, and health information) stored in Amazon S3, and alerts when that data is at risk.	Amazon Inspector: Macie focuses on discovering sensitive data stored in S3 buckets; Inspector focuses on scanning compute resources for software vulnerabilities.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
AWS Secrets Manager	Security	Domain 5	Low	A service for securely storing, managing, retrieving, and automatically rotating application secrets such as database credentials, API keys, and OAuth tokens.	AWS KMS: Secrets Manager stores application credentials and secrets that code needs to authenticate to services; KMS manages the cryptographic keys used to encrypt and decrypt data.

Management and Governance Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
AWS CloudTrail	Management and Governance	Domain 5	High	Records a complete history of all API calls and actions taken within an AWS account by users, roles, and services, creating an auditable activity log for security and compliance investigations.	Amazon CloudWatch: CloudTrail logs who did what and when (API and account activity for auditing); CloudWatch monitors operational performance metrics, resource health, and application logs.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon CloudWatch	Management and Governance	Domains 4, 5	Medium	A monitoring and observability service that collects metrics, logs, and events from AWS resources and applications, enabling automated alerts and dashboards to detect and respond to operational issues.	AWS CloudTrail: CloudWatch monitors what is happening operationally (performance and health); CloudTrail records what actions were taken by whom (activity and change history).
AWS Config	Management and Governance	Domain 5	High	A service that continuously records and tracks the configuration state of AWS resources over time, enabling compliance evaluation against rules and identification of when and how configurations changed.	AWS CloudTrail: Config tracks the state of resource configurations and whether they comply with rules; CloudTrail tracks the API actions that caused those configuration changes to happen.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
AWS Trusted Advisor	Management and Governance	Domain 5	Medium	An automated tool that analyzes an AWS environment across five pillars and provides real-time, actionable recommendations for cost savings, security improvements, fault tolerance, performance, and service limits.	AWS Well-Architected Tool: Trusted Advisor provides continuous automated checks with immediate recommendations; Well-Architected Tool supports structured human-led architecture reviews against AWS best practices.
AWS Well-Architected Tool	Management and Governance	Domain 5	Low	A guided review tool that helps organizations evaluate their workloads against AWS best practices across six pillars: operational excellence, security, reliability, performance efficiency, cost optimization, and sustainability.	AWS Trusted Advisor: Well-Architected Tool is a structured, question-driven review process conducted by architects; Trusted Advisor provides automated, continuous checks without requiring a formal review process.

Developer Tools Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Q	Developer Tools	Domains 1, 2	High	An AI-powered assistant from AWS designed to help developers, IT professionals, and business users with coding, troubleshooting, answering questions about AWS services, and querying enterprise data.	Amazon Lex: Amazon Q is an AI assistant for employees and developers working within AWS tools and enterprise systems; Lex is a platform for building custom-branded conversational bots for end user applications.
Kiro	Developer Tools	Domains 1, 2	Medium	An AI-native software development environment (IDE) from AWS that uses AI agents to help developers design, plan, write, test, and refine code and application architecture more efficiently.	Amazon Q Developer: Kiro is a full AI-native IDE designed around agentic workflows; Amazon Q Developer is an AI coding assistant that works within existing IDEs like VS Code and JetBrains.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Strands Agents	Developer Tools	Domains 2, 3	Medium	An open-source AWS framework that provides building blocks for creating AI agents capable of using tools, managing memory, accessing data, and completing multi-step tasks autonomously.	Amazon Bedrock AgentCore: Strands Agents is the open-source development framework for creating agents; AgentCore is the managed AWS service for deploying, running, and securing those agents in production.

Database Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Aurora	Database	Domain 3	High	A MySQL and PostgreSQL-compatible relational database built for the cloud with high performance and availability, and support for pgvector enabling vector embedding storage for RAG applications.	Amazon RDS: Aurora is a high-performance AWS-built database engine hosted within RDS; RDS is the managed relational database service that hosts Aurora along with MySQL, PostgreSQL, Oracle, and SQL Server.
Amazon DocumentDB	Database	Domain 3	Low	A fully managed document database compatible with MongoDB that stores flexible JSON-like documents, used for application data with varying structures such as user profiles or product catalogs.	Amazon DynamoDB: DocumentDB uses a MongoDB-compatible document model suited to complex document structures; DynamoDB is a key-value and simple document NoSQL database optimized for single-digit millisecond lookups.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon DynamoDB	Database	Domain 3	Low	A fully managed, serverless NoSQL key-value and document database designed for single-digit millisecond response times at any scale, used in applications requiring fast, high-volume lookups.	Amazon RDS: DynamoDB is a NoSQL database optimized for speed and scale with simple access patterns; RDS is a relational database optimized for complex queries and transactional integrity.
Amazon ElastiCache	Database	Domain 3	Low	A fully managed in-memory caching service (Redis or Memcached compatible) used to store frequently accessed data in memory for ultra-low latency retrieval, improving application response times.	Amazon DynamoDB: ElastiCache provides temporary in-memory caching of hot data for microsecond access; DynamoDB is a durable persistent database for storing records with millisecond access.

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon Neptune	Database	Domain 3	Medium	A fully managed graph database service optimized for highly connected data (like social networks or knowledge graphs), and which also supports vector search for AI-powered applications.	Amazon DocumentDB: Neptune is a graph database designed for traversing complex relationships between entities; DocumentDB is a document database for storing and querying flexible document structures.
Amazon RDS	Database	Domain 3	Medium	A managed relational database service that supports multiple engines (MySQL, PostgreSQL, Oracle, SQL Server, MariaDB), with PostgreSQL's pgvector extension enabling vector storage for AI applications.	Amazon Aurora: RDS is the broader managed relational database service that includes many engine options; Aurora is AWS's own high-performance engine option within RDS.

Compute Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon EC2	Compute	Domains 1, 3	Low	The foundational AWS virtual machine service, used to run AI/ML workloads that need dedicated compute - including GPU-based instances for model training and self-hosted model inference.	AWS Lambda: EC2 provides persistent virtual machines suitable for sustained workloads; Lambda runs short-lived functions in response to events without requiring server management.
AWS Lambda	Compute	Domains 1, 3	Low	A serverless compute service that runs code in response to events without requiring server management, commonly used for lightweight AI inference triggers and glue logic in ML pipelines.	Amazon EC2: Lambda is event-driven and best for short, bursty functions; EC2 provides always-available virtual machines for sustained, long-running compute work.

Networking and Content Delivery Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon CloudFront	Networking	Domain 5	Low	A content delivery network (CDN) that distributes web content, API responses, and AI application outputs to users globally through edge locations for reduced latency.	Amazon VPC: CloudFront delivers content to end users at the network edge; VPC is the private internal network where backend AWS resources like databases and servers reside.
Amazon VPC	Networking	Domain 5	Low	A logically isolated private virtual network within AWS where customers deploy resources with full control over IP addressing, subnets, routing, and security - keeping AI workloads off the public internet.	AWS PrivateLink: VPC is the overall private network environment; PrivateLink is a specific feature enabling private connectivity between VPCs and AWS services without traffic ever crossing the public internet.

Storage Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon S3	Storage	Domains 1, 3, 4, 5	High	The foundational AWS object storage service used to store training datasets, model artifacts, logs, documents for RAG knowledge bases, and any other files needed throughout the AI/ML lifecycle.	Amazon S3 Glacier: S3 Standard is for frequently accessed data needing immediate retrieval; S3 Glacier is for archival data that is rarely accessed and can tolerate retrieval delays.
Amazon S3 Glacier	Storage	Domain 5	Low	A low-cost archival storage tier within Amazon S3 designed for data that must be retained for compliance or long-term reference but is rarely if ever accessed.	Amazon S3 Standard: Glacier trades immediate availability for significantly lower storage cost; S3 Standard is priced for data that must be retrieved quickly and frequently.

Container Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
Amazon ECS	Containers	Domain 3	Low	A fully managed container orchestration service that runs and scales Docker containers in AWS without requiring customers to manage the underlying cluster infrastructure.	Amazon EKS: ECS is AWS's own native container management service with a simpler learning curve; EKS is the managed Kubernetes service for organizations that require Kubernetes-compatible tooling.
Amazon EKS	Containers	Domain 3	Low	A managed Kubernetes service that runs Kubernetes container workloads in AWS without requiring installation or management of the Kubernetes control plane.	Amazon ECS: EKS is for Kubernetes workloads requiring portability or Kubernetes-native tooling; ECS is AWS's proprietary service that is simpler to configure for AWS-first environments.

Cloud Financial Management Services

Service Name	Category	Primary Domain(s)	Exam Priority	One-Line Purpose	Commonly Confused With
AWS Budgets	Cloud Financial Management	Domains 2, 5	Low	A service for setting custom cost and usage thresholds for AWS spending, with automated alerts when actual or forecasted costs exceed defined budgets - relevant for managing GenAI token-based pricing costs.	AWS Cost Explorer: Budgets sets spending thresholds and sends alerts when they are approached or exceeded; Cost Explorer provides historical analysis and visualizations of past spending.
AWS Cost Explorer	Cloud Financial Management	Domains 2, 5	Low	A tool for visualizing, analyzing, and understanding AWS costs and usage over time, helping organizations identify patterns, anomalies, and savings opportunities across services including AI/ML services.	AWS Budgets: Cost Explorer is used to analyze and understand historical costs; Budgets is used to set forward-looking spending limits and get notified when they are at risk of being exceeded.

SECTION 5: MASTER TERM LIST

Category 1: AI and ML Fundamentals

Term	Category	Domain	Priority
Artificial Intelligence (AI)	AI and ML Fundamentals	Domain 1	High
Machine Learning (ML)	AI and ML Fundamentals	Domain 1	High
Deep Learning	AI and ML Fundamentals	Domain 1	High
Neural Network	AI and ML Fundamentals	Domain 1	High
Natural Language Processing (NLP)	AI and ML Fundamentals	Domain 1	High
Computer Vision	AI and ML Fundamentals	Domain 1	High
Algorithm	AI and ML Fundamentals	Domain 1	High
Model (ML)	AI and ML Fundamentals	Domain 1	High
Training (ML)	AI and ML Fundamentals	Domain 1	High
Inferencing	AI and ML Fundamentals	Domain 1	High
Supervised Learning	AI and ML Fundamentals	Domain 1	High
Unsupervised Learning	AI and ML Fundamentals	Domain 1	High
Reinforcement Learning	AI and ML Fundamentals	Domain 1	High
Classification	AI and ML Fundamentals	Domain 1	High
Regression	AI and ML Fundamentals	Domain 1	High
Clustering	AI and ML Fundamentals	Domain 1	High
Labeled Data	AI and ML Fundamentals	Domain 1	High
Unlabeled Data	AI and ML Fundamentals	Domain 1	Medium
Feature Engineering	AI and ML Fundamentals	Domain 1	Medium
Overfitting	AI and ML Fundamentals	Domains 1, 4	High
Underfitting	AI and ML Fundamentals	Domains 1, 4	High

Term	Category	Domain	Priority
Bias (Statistical / ML)	AI and ML Fundamentals	Domains 1, 4	High
Variance (ML)	AI and ML Fundamentals	Domains 1, 4	High
Accuracy	AI and ML Fundamentals	Domains 1, 3	High
Precision	AI and ML Fundamentals	Domains 1, 3	High
Recall	AI and ML Fundamentals	Domains 1, 3	High
F1 Score	AI and ML Fundamentals	Domains 1, 3	High
Tabular Data	AI and ML Fundamentals	Domain 1	Medium
Time-Series Data	AI and ML Fundamentals	Domain 1	Medium
Structured Data	AI and ML Fundamentals	Domain 1	Medium
Unstructured Data	AI and ML Fundamentals	Domain 1	Medium
Batch Inferencing	AI and ML Fundamentals	Domain 1	High
Real-Time Inferencing	AI and ML Fundamentals	Domain 1	High
Asynchronous Inferencing	AI and ML Fundamentals	Domain 1	Medium
Serverless Inferencing	AI and ML Fundamentals	Domain 1	Medium
MLOps	AI and ML Fundamentals	Domain 1	Medium
Hyperparameter	AI and ML Fundamentals	Domain 1	Medium
Model Pipeline	AI and ML Fundamentals	Domain 1	Medium
Model Monitoring	AI and ML Fundamentals	Domains 1, 4	Medium
Return on Investment (ROI) for AI	AI and ML Fundamentals	Domain 1	Medium

Category 2: Generative AI and Foundation Models

Term	Category	Domain	Priority
Generative AI (GenAI)	Generative AI and Foundation Models	Domains 1, 2	High

Term	Category	Domain	Priority
Foundation Model (FM)	Generative AI and Foundation Models	Domains 2, 3	High
Large Language Model (LLM)	Generative AI and Foundation Models	Domains 1, 2	High
Token	Generative AI and Foundation Models	Domains 2, 3	High
Tokenization	Generative AI and Foundation Models	Domain 2	Medium
Chunking	Generative AI and Foundation Models	Domains 2, 3	High
Embedding	Generative AI and Foundation Models	Domains 2, 3	High
Vector	Generative AI and Foundation Models	Domains 2, 3	High
Vector Database	Generative AI and Foundation Models	Domain 3	High
Context Window	Generative AI and Foundation Models	Domains 2, 3	High
Context Engineering	Generative AI and Foundation Models	Domain 2	High
Prompt	Generative AI and Foundation Models	Domains 2, 3	High
Transformer Architecture	Generative AI and Foundation Models	Domain 2	Medium
Multi-Modal Model	Generative AI and Foundation Models	Domains 2, 3	High
Diffusion Model	Generative AI and Foundation Models	Domain 2	Medium
Hallucination	Generative AI and Foundation Models	Domains 2, 4, 5	High
Nondeterminism	Generative AI and Foundation Models	Domain 2	Medium
Temperature (Inference Parameter)	Generative AI and Foundation Models	Domain 3	High
Agentic AI	Generative AI and Foundation Models	Domains 1, 2, 3	High
AI Agent	Generative AI and Foundation Models	Domains 2, 3	High
Multi-Agent System	Generative AI and Foundation Models	Domain 2	High
Model Context Protocol (MCP)	Generative AI and Foundation Models	Domain 2	High
Memory Management (Agentic AI)	Generative AI and Foundation Models	Domain 2	Medium
Tool Usage (Agentic AI)	Generative AI and Foundation Models	Domain 2	Medium

Term	Category	Domain	Priority
Workflow Orchestration	Generative AI and Foundation Models	Domains 2, 3	Medium
Token-Based Pricing	Generative AI and Foundation Models	Domains 2, 3	High
FM Lifecycle	Generative AI and Foundation Models	Domain 2	Medium
Provisioned Throughput	Generative AI and Foundation Models	Domains 2, 3	Medium

Category 3: Foundation Model Applications and Techniques

Term	Category	Domain	Priority
Prompt Engineering	Foundation Model Applications	Domains 2, 3	High
Zero-Shot Prompting	Foundation Model Applications	Domain 3	High
Single-Shot Prompting (One-Shot)	Foundation Model Applications	Domain 3	High
Few-Shot Prompting	Foundation Model Applications	Domain 3	High
Chain-of-Thought Prompting	Foundation Model Applications	Domain 3	High
Prompt Template	Foundation Model Applications	Domain 3	Medium
Negative Prompt	Foundation Model Applications	Domain 3	Medium
Retrieval Augmented Generation (RAG)	Foundation Model Applications	Domains 3, 5	High
Fine-Tuning	Foundation Model Applications	Domains 2, 3	High
Instruction Tuning	Foundation Model Applications	Domain 3	Medium
Transfer Learning	Foundation Model Applications	Domain 3	Medium
Continuous Pre-Training	Foundation Model Applications	Domain 3	Medium
Model Distillation	Foundation Model Applications	Domain 3	High
In-Context Learning	Foundation Model Applications	Domain 3	High
Reinforcement Learning from Human Feedback (RLHF)	Foundation Model Applications	Domain 3	High
Grounding	Foundation Model Applications	Domains 3, 5	High

Term	Category	Domain	Priority
ROUGE Score	Foundation Model Applications	Domain 3	High
BLEU Score	Foundation Model Applications	Domain 3	High
BERTScore	Foundation Model Applications	Domain 3	Medium
LLM-as-a-Judge	Foundation Model Applications	Domain 3	High
Knowledge Base (AI)	Foundation Model Applications	Domains 1, 3	High
Inference Parameter	Foundation Model Applications	Domain 3	Medium
Prompt Versioning	Foundation Model Applications	Domain 3	Medium
Human-in-the-Loop Evaluation	Foundation Model Applications	Domains 3, 4	High
Benchmark Dataset	Foundation Model Applications	Domain 3	Medium

Category 4: Responsible AI

Term	Category	Domain	Priority
Responsible AI	Responsible AI	Domain 4	High
Fairness	Responsible AI	Domain 4	High
Inclusivity	Responsible AI	Domain 4	High
Robustness	Responsible AI	Domain 4	Medium
Safety (AI)	Responsible AI	Domains 4, 5	High
Veracity	Responsible AI	Domain 4	Medium
Transparency	Responsible AI	Domain 4	High
Explainability (XAI)	Responsible AI	Domain 4	High
Interpretability	Responsible AI	Domain 4	High
Demographic Bias	Responsible AI	Domain 4	High
Dataset Bias	Responsible AI	Domain 4	High
Model Card	Responsible AI	Domains 4, 5	High
Subgroup Analysis	Responsible AI	Domain 4	Medium
Label Quality	Responsible AI	Domain 4	Medium
Human Audit	Responsible AI	Domain 4	Medium
Environmental Sustainability (AI)	Responsible AI	Domain 4	Medium
Intellectual Property Risk (AI)	Responsible AI	Domain 4	High
Human-Centered Design	Responsible AI	Domain 4	Medium
AI Decision Transparency	Responsible AI	Domain 4	Medium

Term	Category	Domain	Priority
Balanced Dataset	Responsible AI	Domain 4	Medium
Curated Data Source	Responsible AI	Domain 4	Low

Category 5: Security, Compliance, and Governance

Term	Category	Domain	Priority
Shared Responsibility Model	Security and Governance	Domain 5	High
IAM Role	Security and Governance	Domain 5	High
IAM Policy	Security and Governance	Domain 5	High
Encryption at Rest	Security and Governance	Domain 5	High
Encryption in Transit	Security and Governance	Domain 5	High
Data Lineage	Security and Governance	Domain 5	High
Data Cataloging	Security and Governance	Domain 5	Medium
Data Residency	Security and Governance	Domain 5	Medium
Data Governance	Security and Governance	Domain 5	High
Prompt Injection	Security and Governance	Domains 3, 5	High
Prompt Poisoning	Security and Governance	Domains 3, 5	High
Prompt Hijacking	Security and Governance	Domains 3, 5	Medium
Jailbreaking	Security and Governance	Domains 3, 5	Medium
Data Leakage	Security and Governance	Domain 5	High
Output Filtering	Security and Governance	Domain 5	High
Audit Trail	Security and Governance	Domain 5	High
Toxicity Detection	Security and Governance	Domains 4, 5	High
Confidence Scoring	Security and Governance	Domain 5	Medium
Generative AI Security Scoping Matrix	Security and Governance	Domain 5	Medium

Term	Category	Domain	Priority
Governance Framework	Security and Governance	Domain 5	High
Data Lifecycle	Security and Governance	Domain 5	Medium
Privacy-Enhancing Technologies	Security and Governance	Domain 5	Medium
Output Validation	Security and Governance	Domain 5	Medium
AWS PrivateLink	Security and Governance	Domain 5	Medium
Vulnerability Management	Security and Governance	Domain 5	Medium

---

SECTION 6: OFFICIAL TASK STATEMENT SUMMARY

Domain 1: Fundamentals of AI and ML

Task Statement 1.1: Explain basic AI concepts and terminologies

Element	Detail
Task Statement Number	1.1
Full Description	Explain basic AI concepts and terminologies
Cognitive Level	Remember / Understand
Priority	Critical

What the learner must be able to DO:

- *Define* the terms AI, ML, deep learning, neural networks, computer vision, NLP, model, algorithm, training, inferencing, bias, fairness, fit, LLM, GenAI, and agentic AI in plain, non-technical language.
- *Describe* how AI, ML, GenAI, deep learning, and agentic AI are similar to each other and how they differ.
- *Distinguish* between batch, real-time, asynchronous, and serverless inferencing and recognize when each type is most appropriate.
- *Identify* the types of data used in AI models: labeled versus unlabeled data; tabular, time-series, image, and text data; and structured versus unstructured data.
- *Differentiate* between supervised learning, unsupervised learning, and reinforcement learning - including what types of problems each approach is suited for.

Task Statement 1.2: Identify practical use cases for AI

Element	Detail
Task Statement Number	1.2
Full Description	Identify practical use cases for AI
Cognitive Level	Understand / Apply
Priority	Critical

What the learner must be able to DO:

- *Recognize* business scenarios where AI/ML provides clear value: decision support, automation, scalability, and handling tasks that would be impractical for humans alone.
- *Determine* when AI/ML is NOT the right solution - for example, when a rules-based or deterministic answer is required, or when the cost of development and operation outweighs the expected benefit.
- *Select* the appropriate AI/ML technique for a given scenario: regression for predicting continuous numerical values, classification for assigning items to categories, and clustering for discovering natural groupings in unlabeled data.
- *Identify* real-world AI applications and the categories they belong to: computer vision, NLP, speech recognition, recommendation systems, fraud detection, forecasting, knowledge bases, and agentic AI.
- *Explain* the capabilities and appropriate use cases for AWS managed AI/ML services: Amazon SageMaker AI, Amazon Transcribe, Amazon Translate, Amazon Comprehend, Amazon Lex, and Amazon Polly.
- *Identify* when a traditional ML model is more appropriate than a foundation model and vice versa, based on regulatory requirements, explainability needs, cost constraints, and operational considerations.

Task Statement 1.3: Describe the AI/ML development lifecycle

Element	Detail
Task Statement Number	1.3
Full Description	Describe the AI/ML development lifecycle
Cognitive Level	Understand / Apply
Priority	Important

What the learner must be able to DO:

- *Describe* and *differentiate* the components of an AI/ML pipeline, including data collection, preprocessing, feature engineering, model training, evaluation, deployment, and monitoring.
- *Describe* where foundation models come from: open-source pre-trained models and custom-trained models, and the tradeoffs between these sources.
- *Describe* the methods for deploying a model in production: managed API services (where the cloud provider handles infrastructure) and self-hosted APIs (where the customer manages the hosting environment).
- *Identify* relevant AWS services for each stage of an AI/ML pipeline: Amazon Bedrock, Amazon Q, Amazon Quick, Kiro, and SageMaker AI.

- *Explain* core MLOps concepts: experimentation, repeatable processes, scalable systems, managing technical debt, achieving production readiness, model monitoring, and model retraining.
- *Describe* model performance metrics - accuracy, precision, recall, and F1 score - and business metrics - cost per user, development costs, customer feedback, and ROI - and explain what each metric measures and why it matters.

Domain 2: Fundamentals of GenAI

Task Statement 2.1: Explain the basic concepts of generative AI (GenAI)

Element	Detail
Task Statement Number	2.1
Full Description	Explain the basic concepts of generative AI (GenAI)
Cognitive Level	Remember / Understand
Priority	Critical

What the learner must be able to DO:

- *Define* the foundational vocabulary of GenAI: tokens, chunking, embeddings, vectors, prompt engineering, transformer-based LLMs, foundation models (FMs), multi-modal models, and diffusion models.
- *Identify* potential use cases for GenAI models: image, video, and audio generation; text summarization; AI assistants; language translation; code generation; customer service agents; semantic search; and recommendation engines.
- *Describe* the FM lifecycle from start to finish: data selection, model selection, pre-training, fine-tuning, evaluation, deployment, and feedback collection.
- *Describe* how token-based pricing works and its direct effect on cost and performance during inference, including how prompt length and output length affect total token consumption.
- *Describe* the role of context engineering in FM applications - the practice of deliberately designing what information is placed in the model's context window to improve output quality.
- *Define* foundational agentic AI concepts: multi-agent system patterns for complex AI applications, the Model Context Protocol (MCP) and its role in connecting agents to external systems and tools, multi-agent communication patterns, memory management strategies, tool usage, and workflow orchestration.

Task Statement 2.2: Understand the capabilities and limitations of GenAI for solving business problems

Element	Detail
Task Statement Number	2.2
Full Description	Understand the capabilities and limitations of GenAI for solving business problems

Element	Detail
Cognitive Level	Understand / Analyze
Priority	Critical

What the learner must be able to DO:

- *Describe* the advantages of GenAI that make it valuable for business: adaptability across domains, responsiveness to natural language, conversational interaction capabilities, and the ability to generate diverse new content.
- *Identify* the disadvantages and risks of GenAI solutions: hallucinations (confidently incorrect outputs), limited interpretability, potential for inaccuracy, and nondeterminism (the same prompt can produce different outputs each time).
- *Identify* the factors that should inform GenAI model selection for a business problem: model type (text, image, multi-modal), performance requirements, specific capabilities, operational constraints, compliance requirements, cost, latency, and model complexity.
- *Determine* business value and appropriate success metrics for GenAI applications: cross-domain performance, ROI, operational efficiency, conversion rate, average revenue per user, output accuracy, and customer lifetime value.

Task Statement 2.3: Describe AWS infrastructure and technologies for building GenAI applications

Element	Detail
Task Statement Number	2.3
Full Description	Describe AWS infrastructure and technologies for building GenAI applications
Cognitive Level	Remember / Understand
Priority	Critical

What the learner must be able to DO:

- *Identify* the AWS services and features used to develop GenAI applications: Amazon Bedrock, Amazon SageMaker AI, SageMaker JumpStart, Amazon Quick, Kiro, Strands Agents, and Amazon Bedrock AgentCore.
- *Describe* the advantages of using AWS GenAI services rather than building from scratch: accessibility, a lower barrier to entry, operational efficiency, cost-effectiveness, faster time to market, and alignment with business objectives.
- *Describe* the benefits of the AWS infrastructure for GenAI: security certifications, built-in compliance capabilities, clear responsibility boundaries, and safety controls.
- *Describe* the cost tradeoffs of AWS GenAI service options: responsiveness, availability levels, redundancy requirements, performance tiers, regional coverage, token-based pricing structures, provisioned throughput options, and the cost implications of using custom models versus shared models.

Domain 3: Applications of Foundation Models

Task Statement 3.1: Describe design considerations for applications that use foundation models (FMs)

Element	Detail
Task Statement Number	3.1
Full Description	Describe design considerations for applications that use foundation models (FMs)
Cognitive Level	Understand / Apply / Analyze
Priority	Critical

What the learner must be able to DO:

- *Identify* the selection criteria for choosing an FM for a given use case: cost, modality (text-only, image, audio, multi-modal), latency requirements, multi-lingual support, model size, model complexity, available customization options, input and output length limits, and prompt caching availability.
- *Describe* how inference parameters affect model responses, particularly temperature (which controls how creative or conservative outputs are) and input/output length limits.
- *Define* Retrieval Augmented Generation (RAG) and *describe* its business applications - including how Amazon Bedrock Knowledge Bases enables organizations to ground model responses in their own company data.
- *Identify* the AWS services that can store vector embeddings for use in RAG systems: Amazon OpenSearch Service, Amazon Aurora, Amazon Neptune, and Amazon RDS for PostgreSQL (with the pgvector extension).
- *Explain* the cost and complexity tradeoffs between different FM customization approaches: pre-training a model from scratch (highest cost), fine-tuning an existing model, in-context learning, RAG, and model distillation (creating a smaller, faster model trained on outputs from a larger model).
- *Define* the role of AI agents and *describe* practical business applications where agents provide value through autonomous multi-step task completion.

Task Statement 3.2: Choose effective prompt engineering techniques

Element	Detail
Task Statement Number	3.2
Full Description	Choose effective prompt engineering techniques
Cognitive Level	Apply / Analyze
Priority	Critical

What the learner must be able to DO:

- *Define* the core concepts and components of a well-structured prompt: context (background information the model needs), instruction (the specific task to perform), and negative prompts (explicit guidance on what the model should avoid).

- *Define* and *differentiate* prompt engineering techniques: chain-of-thought prompting (guiding the model to reason step by step), zero-shot prompting (asking with no examples), single-shot prompting (providing one example), few-shot prompting (providing multiple examples), and prompt templates (reusable structured prompt formats).
- *Identify* and *describe* the benefits and best practices of effective prompt engineering: improving response quality, supporting systematic experimentation, using guardrails to constrain outputs, achieving appropriate specificity and concision, and using multiple instructions for complex tasks.
- *Define* the risks and attack surfaces of prompt engineering: prompt exposure (leaking system prompt content), prompt poisoning (injecting malicious instructions through data), prompt hijacking (overriding intended behavior), and jailbreaking (bypassing safety restrictions).
- *Describe* prompt versioning and management strategies using Amazon Bedrock Prompt Management, including why version control for prompts matters in production AI applications.

Task Statement 3.3: Describe the training and fine-tuning process for FMs

Element	Detail
Task Statement Number	3.3
Full Description	Describe the training and fine-tuning process for FMs
Cognitive Level	Understand / Apply
Priority	Important

What the learner must be able to DO:

- *Describe* the key stages of training a foundation model: pre-training (learning from a massive general dataset), fine-tuning (adapting to a specific task or domain), continuous pre-training (updating with new data over time), and distillation (compressing knowledge into a smaller, faster model).
- *Define* the methods available for fine-tuning an FM: instruction tuning (training on task-instruction pairs), domain-specific adaptation (training on industry-specific data), transfer learning (applying knowledge from one context to another), and continuous pre-training (ongoing updates without full retraining).
- *Describe* how to prepare data for fine-tuning: data curation (selecting relevant, high-quality examples), governance (ensuring the data is legally usable and properly licensed), sizing (ensuring sufficient training examples), labeling (accurately annotating training data), ensuring representativeness (covering the full range of expected inputs), and using Reinforcement Learning from Human Feedback (RLHF) to incorporate human quality judgments into the training process.

Task Statement 3.4: Describe methods to evaluate FM performance

Element	Detail
Task Statement Number	3.4
Full Description	Describe methods to evaluate FM performance
Cognitive Level	Apply / Analyze
Priority	Critical

What the learner must be able to DO:

- *Determine* which evaluation approaches are appropriate for a given situation: human-in-the-loop evaluation (having people rate model outputs), benchmark datasets (comparing outputs against known reference answers), and Amazon Bedrock Model Evaluation (AWS's managed evaluation service).
- *Identify* automated metrics for assessing FM quality: ROUGE (Recall-Oriented Understudy for Gisting Evaluation, used for summarization quality), BLEU (Bilingual Evaluation Understudy, used for translation quality), BERTScore (semantic similarity between generated and reference text), and LLM-as-a-judge (using a second, trusted LLM to evaluate the quality of another model's outputs).
- *Determine* whether an FM is meeting business objectives by evaluating outcomes like productivity improvements, user engagement levels, and task completion rates.
- *Identify* appropriate evaluation approaches for applications built on top of FMs: assessing RAG pipeline quality (relevance and groundedness of retrieved content), evaluating agent performance (task success rate and efficiency), and evaluating automated workflows (accuracy and reliability).
- *Identify* business objective alignment metrics for AI applications: task completion rate (the percentage of user goals successfully achieved), user satisfaction scores, and cost per interaction (the total cost to handle one user request end to end).

Domain 4: Guidelines for Responsible AI

Task Statement 4.1: Explain the development of AI systems that are responsible

Element	Detail
Task Statement Number	4.1
Full Description	Explain the development of AI systems that are responsible
Cognitive Level	Understand / Apply / Analyze
Priority	Critical

What the learner must be able to DO:

- *Identify* the core features of responsible AI: bias (systematic errors that produce unfair results), fairness (equal and just treatment of all groups), inclusivity (designing AI to work well for diverse populations), robustness (reliable performance under varied conditions), safety (avoiding harmful outputs or actions), and veracity (producing truthful and accurate outputs).

- *Explain* how to use Amazon Bedrock Guardrails to identify and enforce responsible AI features - including content filters, topic restrictions, and grounding checks.
- *Define* responsible practices for model selection, including considering the environmental impact and energy consumption of large AI models, and the sustainability implications of training and deployment choices.
- *Identify* the legal risks associated with GenAI: intellectual property infringement from training data or generated outputs, biased model outputs that could create legal liability, loss of customer trust from AI failures, end user risks from relying on AI for high-stakes decisions, and hallucination-based misinformation.
- *Identify* characteristics of datasets that support responsible AI: inclusivity (representing diverse populations), diversity (covering varied scenarios and edge cases), use of curated and vetted data sources, and balanced class representation to prevent skewed learning.
- *Describe* the effects of bias and variance on model fairness and performance: how statistical bias can disproportionately affect demographic groups, cause systematic inaccuracy, and lead to overfitting or underfitting.
- *Describe* the AWS tools used to detect and monitor bias, trustworthiness, and truthfulness in AI systems: analyzing label quality in training datasets, conducting human audits, performing subgroup analysis, using Amazon SageMaker Clarify for bias detection, SageMaker Model Monitor for ongoing performance tracking, and Amazon Augmented AI (A2I) for incorporating human review.

Task Statement 4.2: Recognize the importance of transparent and explainable models

Element	Detail
Task Statement Number	4.2
Full Description	Recognize the importance of transparent and explainable models
Cognitive Level	Understand / Analyze
Priority	Important

What the learner must be able to DO:

- *Describe* the meaningful differences between AI models that are transparent and explainable versus those that are not - including why this difference matters for trust, accountability, regulatory compliance, and the ability to identify and correct errors.
- *Describe* the AWS tools and practices that support transparency and explainability: Amazon SageMaker Model Cards (structured documentation of model behavior and limitations), SageMaker Clarify (feature importance and bias explanations), Amazon Bedrock Model Evaluations (quantitative model assessment), and the use of open-source models and transparent data and licensing documentation.
- *Identify* the inherent tradeoffs between model safety and transparency - for example, that highly performant deep learning models tend to be harder to explain than simpler, more interpretable models, and that some explanations may reveal sensitive model details.
- *Describe* the principles of human-centered design for explainable AI: designing user-facing systems that provide meaningful explanations of AI decisions, implementing user feedback

mechanisms so people can report when AI outputs seem wrong or unfair, and ensuring AI decision transparency so users understand what role AI played in any outcome that affects them.

Domain 5: Security, Compliance, and Governance for AI Solutions

Task Statement 5.1: Explain methods to secure AI systems

Element	Detail
Task Statement Number	5.1
Full Description	Explain methods to secure AI systems
Cognitive Level	Understand / Apply
Priority	Critical

What the learner must be able to DO:

- *Identify* AWS services and features used to secure AI systems: IAM roles, policies, and permissions (access control); encryption using AWS KMS (protecting data at rest and in transit); Amazon Macie (detecting sensitive data in S3); AWS PrivateLink (private network connectivity); the AWS shared responsibility model (clarifying who is responsible for what); Amazon Bedrock AgentCore Identity and AgentCore Policy (securing AI agents); and Amazon Bedrock Guardrails (filtering and controlling model outputs).
- *Describe* source citation and data origin documentation practices: data lineage (tracking where data came from and how it was transformed), data cataloging (maintaining a structured inventory of datasets), and using Amazon SageMaker Model Cards to document model provenance and training data sources.
- *Describe* best practices for secure data engineering in AI contexts: assessing data quality before use, implementing privacy-enhancing technologies (such as differential privacy and data masking), enforcing strict data access controls, and ensuring data integrity throughout the AI/ML pipeline.
- *Describe* the security and privacy considerations specific to AI systems: application security, threat detection, vulnerability management, infrastructure protection, prompt injection attacks, encryption at rest and in transit, data leakage prevention, output filtering and validation, audit trail and logging requirements for AI interactions, and toxicity detection in AI outputs.
- *Describe* hallucination detection methods and grounding techniques that improve AI output accuracy: RAG grounding (anchoring responses in retrieved, verified sources), output validation (checking model outputs against defined criteria), and confidence scoring (providing a measure of how certain the model is about its output).

Task Statement 5.2: Recognize governance and compliance regulations for AI systems

Element	Detail
Task Statement Number	5.2
Full Description	Recognize governance and compliance regulations for AI systems
Cognitive Level	Remember / Understand / Apply
Priority	Important

What the learner must be able to DO:

- *Identify* the AWS services that support governance and regulatory compliance for AI: AWS Config (configuration compliance tracking), Amazon Inspector (vulnerability scanning), AWS Audit Manager (automated evidence collection), AWS Artifact (access to AWS compliance reports), AWS CloudTrail (API activity logging), and AWS Trusted Advisor (best practice recommendations).
- *Describe* data governance strategies relevant to AI workloads: managing data lifecycles (creation through deletion), logging data access and usage, enforcing data residency requirements (keeping data in specified geographic regions), monitoring and observing data use, and defining data retention policies.
- *Describe* the processes organizations should follow to maintain AI governance: establishing and enforcing policies, defining regular review cadences, implementing systematic review strategies, using recognized governance frameworks such as the Generative AI Security Scoping Matrix, meeting transparency standards for AI use, and fulfilling team training requirements to ensure human oversight.

SECTION 7: RECOMMENDED LEARNER STUDY SEQUENCE

The sequence below is designed for a learner starting with limited familiarity with AI/ML concepts. It moves from orientation through foundational vocabulary, domain-by-domain learning, reference review, and then active practice and assessment. Each step builds on the previous one. Estimated times assume focused, active study - not passive reading.

Step 1: Task 2 - Learner Orientation and Program Overview Guide

Why here: Before engaging with any content, learners need to understand the big picture: what the AIF-C01 exam tests, what is expected of them, and how this program is organized. Setting context first reduces uncertainty and improves retention of everything that follows. This is the entry point for every new learner.

After completing this step, learners will be able to: Describe what the AIF-C01 exam tests at a high level, understand the five domain structure, know what to expect throughout the study program, confirm whether they have the recommended baseline AWS knowledge, and create a personal study schedule.

Estimated time: 30-45 minutes.

Step 2: Task 1 (Sections 2 and 4 only) - Master Project Blueprint Reference Read

Why here: Before studying any domain in depth, learners benefit from reading the Domain Summary Table (Section 2) and the In-Scope AWS Services Registry (Section 4) of this document. This creates a mental map of the subject matter, surfaces the most important AWS services immediately, and helps learners understand how the five domains connect to each other.

After completing this step, learners will be able to: Name all five domains and their approximate exam weights, identify the most important AWS services for the exam by priority level, and understand which domains deserve the greatest study focus.

Estimated time: 45-60 minutes.

Step 3: Task 3 - Domain 1 Study Module: Fundamentals of AI and ML

Why here: Domain 1 is the essential foundation for the entire certification. The vocabulary and concepts introduced here - AI, ML, supervised learning, classification, regression, the ML pipeline - appear in every other domain. Learners must establish this vocabulary before moving forward. Without Domain 1, Domains 2 through 5 will feel disconnected and hard to follow.

After completing this step, learners will be able to: Define core AI/ML terms, distinguish between AI, ML, and GenAI, describe the types of learning and inferencing, identify appropriate ML techniques for business scenarios, name the AWS managed AI services and what they do, and describe the stages of an AI/ML development pipeline.

Estimated time: 3-4 hours.

Step 4: Task 4 - Domain 2 Study Module: Fundamentals of GenAI

Why here: Domain 2 builds directly on the Domain 1 vocabulary and introduces the world of generative AI. It also introduces Amazon Bedrock, which is central to Domains 3, 4, and 5. With 12 scored questions (~24%), this is the second most heavily tested domain, making thorough understanding here essential.

After completing this step, learners will be able to: Explain what tokens, embeddings, and context windows are in plain language; describe how foundation models work at a conceptual level; identify the advantages and limitations of GenAI for business use; name the AWS services used to build GenAI applications; and describe the core concepts of agentic AI including the Model Context Protocol.

Estimated time: 3-4 hours.

Step 5: Task 5 - Domain 3 Study Module: Applications of Foundation Models

Why here: Domain 3 is the most heavily weighted domain at 28% (approximately 14 scored questions). It requires applying the concepts from Domains 1 and 2, rather than just recalling them. It covers RAG, prompt engineering, fine-tuning, and performance evaluation - all topics where learners are expected to make judgment calls, not just recognize definitions. This module requires the most study time of any in the program.

After completing this step, learners will be able to: Choose the right foundation model for a given scenario based on defined criteria, write effective prompts using standard techniques, explain what RAG is and why it matters for grounding AI in company data, describe the fine-tuning process and when to use it, and identify the metrics used to evaluate model performance against both technical and business objectives.

Estimated time: 4-5 hours.

Step 6: Task 6 - Domain 4 Study Module: Guidelines for Responsible AI

Why here: Domain 4 concepts are closely connected to Domain 3 (responsible prompt engineering, evaluating for bias) and Domain 5 (data security and governance). Studying it after Domain 3 allows learners to connect responsible AI principles to the foundation model use cases they already understand. The AWS tools in this domain - especially SageMaker Clarify and Bedrock Guardrails - appear across multiple domains.

After completing this step, learners will be able to: Identify the six features of responsible AI (bias, fairness, inclusivity, robustness, safety, veracity), name the AWS tools used to detect and address bias in AI systems, explain the importance of model transparency and explainability, describe the legal risks of GenAI deployment, and identify the characteristics of a responsible training dataset.

Estimated time: 2-3 hours.

Step 7: Task 7 - Domain 5 Study Module: Security, Compliance, and Governance for AI Solutions

Why here: Domain 5 synthesizes AWS security knowledge from general AWS experience and applies it to AI-specific risks and governance requirements. Topics like prompt injection, data governance, and AI audit trails are best understood after learners already grasp how AI systems work from Domains 1-4. This domain also reinforces the AWS Shared Responsibility Model in the AI context.

After completing this step, learners will be able to: Identify the AWS security services relevant to AI workloads and what each one does, describe how the AWS Shared Responsibility Model applies to AI, explain prompt injection and data leakage as AI-specific security threats, describe hallucination grounding techniques, and outline the governance processes organizations should follow when deploying AI.

Estimated time: 2-3 hours.

Step 8: Task 9 - Master Glossary

Why here: After completing all five domain study modules, learners have encountered every key term at least once in context. Reviewing the Master Glossary at this stage serves as reinforcement - filling in gaps, sharpening definitions, and connecting terms across domains. This is active review, not passive reading; learners should self-quiz rather than simply scan the list.

After completing this step, learners will be able to: Quickly and accurately define any key term from any domain without referring back to the study modules, and recognize terms in the phrasing used by exam questions.

Estimated time: 1-2 hours.

Step 9: Task 8 - AWS Services Quick Reference Guide

Why here: After studying all five domains, learners have encountered every in-scope AWS service in the context of learning objectives. A focused, deliberate review of the services guide at this point reinforces which services do what, clarifies the differences between commonly confused services, and builds the confidence needed for service-selection questions on the exam.

After completing this step, learners will be able to: State the primary purpose of every in-scope AWS service, clearly differentiate between commonly confused service pairs (such as Bedrock vs. SageMaker AI, Macie vs. Inspector, CloudTrail vs. CloudWatch), and select the correct service for a given exam scenario.

Estimated time: 1-2 hours.

Step 10: Task 10 - Visual Cheat Sheet and One-Page Summary

Why here: At this point in the program, learners have studied all content and reviewed key terms and services. The cheat sheet provides a compressed, high-density view of the most critical facts from all five domains. Using it before starting practice questions helps learners identify which facts they know well and which still need reinforcement.

After completing this step, learners will be able to: Self-test rapidly on the most critical facts, terms, service names, and domain relationships across all five domains, and identify any remaining knowledge gaps before moving into formal practice questions.

Estimated time: 30-45 minutes.

Step 11: Task 11 - Master Flashcard Deck

Why here: Active recall through flashcards is one of the most effective evidence-based study techniques. After completing all study modules and reference materials, learners are ready to test their memory systematically rather than just rereading content. Flashcards build the speed and automatic recall that exam conditions demand.

After completing this step, learners will be able to: Rapidly recall definitions, service purposes, and key distinctions without needing to think through each concept from scratch, building the speed and confidence needed for 65 questions in 90 minutes.

Estimated time: 2-4 hours spread across multiple sessions (spaced repetition produces the best retention results).

Step 12: Task 12 - Domain Practice Question Banks

Why here: Domain-specific practice questions apply knowledge in the same format as actual exam questions. Working through questions domain by domain allows learners to pinpoint their weakest areas before attempting a full-length simulated exam. When an answer is wrong, learners should return to the corresponding study module section and review before moving on.

After completing this step, learners will be able to: Identify their strongest and weakest domains by name, recognize how exam questions are typically worded and structured across all four question types, understand which service or concept a given type of question is likely testing, and feel confident about exam question format before the timed full-length exam.

Estimated time: 2-3 hours (including reviewing incorrect answers thoroughly).

Step 13: Task 13 - Full-Length Practice Exam (65 Questions)

Why here: After targeted domain-level practice, learners are ready to experience a complete simulated exam under conditions that mirror the real thing. Taking 65 questions under a 90-minute time limit builds exam stamina, reveals whether study efforts have been comprehensive, and provides a scaled score to compare against the 700-point passing threshold.

After completing this step, learners will be able to: Assess their overall exam readiness, manage time appropriately across 65 questions, experience the complete range of question types in a mixed-domain context, and identify any remaining specific topics that need targeted last-minute review.

Estimated time: 90 minutes for the exam itself; additional review time after.

Step 14: Task 14 - Practice Exam Answer Key and Detailed Rationale Guide

Why here: The rationale guide is the highest-value document for learners who have already completed the practice exam. It explains not just which answer is correct, but why - and why every incorrect option fails. Studying wrong answers with explanation is one of the most efficient learning activities in exam preparation.

After completing this step, learners will be able to: Understand the reasoning behind correct answers for every question type, recognize the distractor patterns that exam questions use to mislead, identify and close any remaining knowledge gaps with targeted review of specific study module sections, and approach the real exam with well-founded confidence.

Estimated time: 1-2 hours.

Step 15: Task 15 - Exam Day Strategy Guide and Final Readiness Checklist

Why here: The final step in preparation is reviewing strategy and logistics, not adding new content. At this point, learners know the material. This guide ensures they know how to perform well under exam conditions - how to pace themselves, approach different question types, handle uncertainty, and arrive ready on exam day.

After completing this step, learners will be able to: Manage time effectively across 90 minutes, apply specific techniques for multiple choice, multiple response, ordering, and matching question types, confidently handle questions they are uncertain about, confirm they are technically and logistically prepared for online proctored or testing center delivery, and walk into the exam calm, organized, and ready.

Estimated time: 30-45 minutes.

End of Master Blueprint - Task 1 of 15 Document version: 1.0 | Based on AIF-C01 Exam Guide Version 1.1 (April 30, 2026) Next task: Task 2 - Learner Orientation and Program Overview Guide