

AWS Certified AI Practitioner (AIF-C01) - Full-Length Practice Exam

Learner Version - For Self-Assessment

Task 10 of 15 / Companion to the Master Project Blueprint Based on AIF-C01 Exam Guide Version 1.1 (April 30, 2026)

Instructions

- This practice exam contains **65 questions**. Recommended time: **90 minutes**.
- **Single-select questions:** choose ONE answer (A, B, C, or D).
- **Multiple-response questions** will say “Select TWO” or “Select THREE.” Read carefully - you must select every correct option and no incorrect options to receive credit. There is no partial credit.
- **Ordering questions** ask you to place a list of items into the correct sequence. All items must be placed correctly to receive credit.
- **Matching questions** ask you to match a list of concepts to a list of examples or definitions. All pairs must be correct to receive credit.
- Do NOT reference study materials, the Master Glossary, or any other program document while taking this exam.
- Mark any question you are unsure about and return to it before finishing.
- Answer every question, even if you are uncertain. Unanswered questions are scored as incorrect, and there is no penalty for guessing.
- Record all of your answers on the Answer Sheet at the end of this document.
- Questions are presented in mixed order across all five domains, exactly as they would appear on the real exam. They are not grouped by domain.

Disclaimer: This is original practice material created for this study program. It is not an official AWS exam and has not been produced, reviewed, or endorsed by Amazon Web Services.

SECTION 2: THE 65 QUESTIONS

Question

1

Place
the
fol-
low-
ing
stages
of
the
foun-
da-
tion
model
(FM)
life-
cycle
in
the
cor-
rect
or-
der,
from
first
to
last.

1. Fine-tuning
2. Data selection
3. Deployment
4. Pre-training
5. Evaluation

*Arrange
items
1-5
into
the
cor-
rect
se-
quence.*

Question
2

A
retail
oper-
a-
tions
man-
ager
de-
scribes
a
new
sys-
tem:
“It
uses
lay-
ers of
inter-
con-
nected
nodes
to
auto-
mati-
cally
learn
pat-
terns
from
thou-
sands
of
his-
tori-
cal
im-
ages
of
dam-
aged
pack-
ages,
with-
out
any-
one
man-
ually
defin-
ing⁴
what
‘dam-

A.
Robotic
pro-
cess
au-
toma-
tion
B.
Deep
learn-
ing C.
Rules-
based
ex-
pert
sys-
tem
D.
Busi-
ness
intel-
li-
gence
dash-
board

Question
3

A
startup
needs
a
foun-
da-
tion
model
that
sup-
ports
both
En-
glish
and
Span-
ish
equally
well,
re-
sponds
in
un-
der
one
sec-
ond,
and
fits
within
a
tight
monthly
bud-
get.
Which
set of
fac-
tors
is the
com-
pany
eval-
uat-
ing?

A.
Model
train-
ing
data
label-
ing
qual-
ity
only

B.
FM
selec-
tion
crite-
ria,
in-
clud-
ing
multi-
lingual
sup-
port,
la-
tency,
and
cost

C.
RLHF
train-
ing
stages

D.
Diffu-
sion
model
archi-
tec-
ture

Question
4

A developer explains that before a large language model can process a long document for a knowledge base, the document is split into smaller pieces, each piece is converted into a list of numbers that captures its meaning, and those numbers

-
- A.
Tok-
eniza-
tion,
then
fine-
tuning
- B.
Chunk-
ing,
then
em-
bed-
ding
- C.
Prompt
engi-
neer-
ing,
then
distil-
la-
tion
- D.
Con-
text
engi-
neer-
ing,
then
pre-
training
-

Question
5

A company wants to ensure that its AI hiring tool produces truthful, accurate outputs and does not generate fabricated claims about a candidate's qualifications. Which core feature of responsible AI is most directly being addressed?

A.
Ve-
rac-
ity B.
Inclu-
sivity
C.
Ro-
bust-
ness
D.
Con-
fi-
dence
scor-
ing

Question
6

A
cre-
ative
writ-
ing
appli-
ca-
tion
wants
its
foun-
da-
tion
model
to
gen-
erate
more
var-
ied,
sur-
pris-
ing,
and
un-
con-
ven-
tional
story
ideas,
even
at
the
risk
of
occa-
sional
less
co-
her-
ent
out-
put.
Which
infer-
ence
pa-
rame-
ter
ad-
just-
ment

-
- A.
De-
crease
the
tem-
pera-
ture
set-
ting
- B. In-
crease
the
tem-
pera-
ture
set-
ting
- C.
De-
crease
the
con-
text
win-
dow
- D. In-
crease
the
num-
ber
of
fine-
tuning
epochs
-

Question
7

A
com-
pany
runs
a
GenAI
appli-
ca-
tion
using
Ama-
zon
Bedrock.
Un-
der
the
AWS
Shared
Re-
spon-
sibil-
ity
Model,
which
of
the
fol-
low-
ing is
the
cus-
tomer's
re-
spon-
sibil-
ity,
not
AWS's?

A.
Se-
cur-
ing
the
phys-
ical
data
cen-
ters
that
host
Bedrock's
in-
fras-
truc-
ture

B.
Con-
figur-
ing
IAM
per-
mis-
sions
and
prop-
erly
man-
aging
the
data
the
cus-
tomer
sends
to
and
stores
within
the
appli-
ca-
tion

C.
Main-
tain-
ing
the
un-
derly-
ing

Question
8

A bank wants to score 2 million loan applications submitted over the past month and generate a single report tonight. There is no requirement for an immediate response to any individual application. Which type of infer-

17

enc-

ing

1

A.
Real-
time
infer-
enc-
ing B.
Batch
infer-
enc-
ing
C.
Asyn-
chronous
infer-
enc-
ing
D.
Server-
less
infer-
enc-
ing

Question
9

A team is building a chatbot using a foundation model and notices that when the conversation gets very long, the model starts to “forget” details mentioned earlier in the same conversation. What concept explains this behavior?

A.
The
model
has
reached
its
con-
text
win-
dow
limit

B.
The
model
is
expe-
rienc-
ing
hallu-
cina-
tion

C.
The
model
has
been
fine-
tuned
incor-
rectly

D.
The
model
is
using
batch
infer-
enc-
ing

Question
10

A
law
firm
wants
its
AI
assis-
tant
to
an-
swer
ques-
tions
using
the
firm's
own
inter-
nal
case
files
and
prece-
dent
docu-
ments,
rather
than
rely-
ing
only
on
the
model's
gen-
eral
train-
ing
data,
so
that
an-
swers
are
grounded
in
the
firm's
ac-
tual
records.
Which

A.
Model
distil-
la-
tion
B.
Re-
trieval
Aug-
mented
Gen-
era-
tion
(RAG)
C.
Prompt
poi-
son-
ing
D.
Rein-
force-
ment
learn-
ing

Question
11

A
com-
pany
wants
to
pre-
vent
its
customer-
facing
GenAI
chat-
bot
from
dis-
cussing
com-
peti-
tor
prod-
ucts
or
pro-
vid-
ing
medi-
cal
ad-
vice,
re-
gard-
less
of
how
users
phrase
their
re-
quests.
Which
AWS
capa-
bility
is de-
signed
for
this
pur-
pose?

A.
Ama-
zon
Sage-
Maker
Clar-
ify B.
Ama-
zon
Bedrock
Guardrails
C.
AWS
Cloud-
Trail
D.
Ama-
zon
In-
spec-
tor

Question
12

A
mar-
ket-
ing
team
is
eval-
uat-
ing
risks
be-
fore
rolling
out a
gen-
era-
tive
AI
tool
to
draft
customer-
facing
emails
with-
out
hu-
man
re-
view.
Se-
lect
TWO
dis-
ad-
van-
tages
of
GenAI
that
are
most
rele-
vant
to
this
deci-
sion.

A.
Hal-
luci-
na-
tion -
the
model
may
gen-
erate
plausible-
sounding
but
factu-
ally
incor-
rect
con-
tent

B.
GenAI
mod-
els
al-
ways
re-
quire
server
in-
fras-
truc-
ture
to be
man-
ually
pro-
vi-
sioned

C.
Non-
de-
ter-
min-
ism -
the
same
prompt
can
pro-
duce
dif-
fer-

*Select
TWO.*

Question
13

A
hos-
pital
is
build-
ing
an
AI
sys-
tem
that
will
ana-
lyze
MRI
scans,
doc-
tors'
hand-
writ-
ten
notes,
and
a
struc-
tured
spread-
sheet
of pa-
tient
lab
val-
ues.
Which
state-
ment
cor-
rectly
clas-
sifies
these
data
types?

A.
All
three
are
struc-
tured
data
be-
cause
they
all
de-
scribe
pa-
tients

B.
MRI
scans
and
hand-
writ-
ten
notes
are
un-
struc-
tured
data;
the
lab
value
spread-
sheet
is
struc-
tured
data

C.
MRI
scans
are
struc-
tured
data;
notes
and
lab
val-
ues
are
un-
struc-

Question

14

Match

each

prompt

engi-

neer-

ing

con-

cept

(1-3)

to

the

ex-

am-

ple

that

best

illus-

trates

it (A-

C).

Concepts:

1.

Con-

text

2. In-

struc-

tion

3.

Nega-

tive

prompt

Examples:

A.
“Do
not
in-
clude
any
pric-
ing
infor-
ma-
tion
in
your
re-
sponse.”

B.
“Sum-
ma-
rize
the
at-
tached
cus-
tomer
com-
plaint
in
two
sen-
tences.”

C.
“The
cus-
tomer
has
been
a
mem-
ber
for
five
years
and
previ-
ously
con-
tacted
sup-
port
twice
about
1.1.1.1

*Match
each
num-
bered
con-
cept
to its
corre-
spond-
ing
let-
tered
ex-
am-
ple.*

Question
15

A
com-
pany
wants
to
auto-
mati-
cally
dis-
cover
and
clas-
sify
any
per-
son-
ally
iden-
tifi-
able
infor-
ma-
tion
(PII)
that
may
have
acci-
den-
tally
been
stored
in an
S3
bucket
used
for
AI
train-
ing
data.
Which
AWS
ser-
vice
should
they
use?

A.
Ama-
zon
In-
spec-
tor B.
Ama-
zon
Ma-
cie C.
AWS
Pri-
vateLink
D.
AWS
Key
Man-
age-
ment
Ser-
vice

Question
16

A
com-
pany
wants
to
use
an
AI
model
that
can
ac-
cept
both
an
im-
age
and
a
text
ques-
tion
about
that
im-
age,
then
gen-
erate
a rel-
evant
text
an-
swer.
Which
type
of
model
is
best
suited
for
this
use
case?

A. A
diffu-
sion
model
B. A
multi-
modal
model
C. A
single-
modality
clus-
ter-
ing
model
D. A
tradi-
tional
re-
gres-
sion
model

Question
17

A
com-
pany
al-
ready
uses
Ama-
zon
RDS
for
Post-
greSQL
for
its
appli-
ca-
tion
data
and
wants
to
add
vec-
tor
stor-
age
for a
RAG
appli-
ca-
tion
with-
out
intro-
duc-
ing
an
en-
tirely
new
database
tech-
nol-
ogy.
Which
ap-
proach
best
fits
this
need?

A.
Mi-
grate
to
Ama-
zon
Nep-
tune
im-
medi-
ately

B.
Use
the
pgvec-
tor
ex-
ten-
sion
within
Ama-
zon
RDS
for
Post-
greSQL
to
store
vec-
tor
em-
bed-
dings

C.
Store
vec-
tor
em-
bed-
dings
as
plain
text
files
in
Ama-
zon
S3
Glacier

38
D.
Use
Ama-

Question

18

Place

the

fol-

low-

ing

stages

of a

typi-

cal

AI/ML

de-

vel-

op-

ment

pipeline

into

the

cor-

rect

or-

der,

from

first

to

last.

1.
Model
train-
ing 2.
Data
col-
lec-
tion
3.
De-
ploy-
ment
4.
Data
pre-
pro-
cess-
ing 5.
Model
eval-
ua-
tion
*Arrange
items
1-5
into
the
cor-
rect
se-
quence.*

Question
19

When
eval-
uat-
ing
two
foun-
da-
tion
mod-
els
with
simi-
lar
accu-
racy,
a
com-
pany's
re-
spon-
sible
AI
pol-
icy
re-
quires
con-
sider-
ing
the
envi-
ron-
men-
tal
foot-
print
of
train-
ing
and
run-
ning
each
model
be-
fore
mak-
ing a
final
41
selec-
tion.
Which

A.
Fair-
ness
B.
Envi-
ron-
men-
tal
sus-
tain-
abil-
ity
con-
sider-
a-
tions
in
model
selec-
tion
C.
Ex-
plain-
abil-
ity D.
Data
lin-
eage

Question
20

A
com-
pany's
chat-
bot
appli-
ca-
tion
sends
very
long
back-
ground
in-
struc-
tions
with
every
sin-
gle
user
mes-
sage
to
main-
tain
con-
sis-
tent
be-
hav-
ior,
in
addi-
tion
to
the
model's
lengthy
re-
sponse
each
time.
The
team
no-
tices
their
monthly
Amazon
Bedrock

A.
Bedrock
charges
a flat
monthly
fee
re-
gard-
less
of us-
age,
so us-
age
pat-
terns
are
irrel-
evant

B.
Both
the
long
input
in-
struc-
tions
and
the
long
out-
put
re-
sponses
count
to-
ward
bill-
able
to-
kens,
so
longer
prompts
and
longer
re-
sponses
di-
rectly
in-
crease
cost

C.

Question
21

A
com-
pany
wants
to
teach
a
foun-
da-
tion
model
company-
specific
ter-
mi-
nol-
ogy
and
tone,
has a
mod-
est
bud-
get,
and
does
not
want
to re-
train
the
en-
tire
model
from
scratch
or
build
a per-
ma-
nent
re-
trieval
sys-
tem.
Which
cus-
tomiza-
tion
ap-
proach
offers

A.
Pre-
training
a
new
model
from
scratch

B.
Fine-
tuning
the
exist-
ing
foun-
da-
tion
model
on
company-
specific
ex-
am-
ples

C.
Do-
ing
noth-
ing
and
using
the
base
model
as-is

D.
Build-
ing a
brand-
new
trans-
former
archi-
tec-
ture

Question
22

An
audi-
tor
wants
to
un-
der-
stand
ex-
actly
where
the
train-
ing
data
for a
de-
ployed
model
came
from
and
how
it
was
trans-
formed
be-
fore
being
used.
Which
con-
cept
and
AWS
tool
com-
bina-
tion
most
di-
rectly
sup-
ports
this
need?

A.
En-
cryp-
tion
in
tran-
sit,
using
AWS
KMS

B.
Data
lin-
eage,
docu-
mented
using
Ama-
zon
Sage-
Maker
Model
Cards

C.
Token-
based
pric-
ing,
tracked
using
AWS
Bud-
gets

D.
Prompt
engi-
neer-
ing,
man-
aged
using
Ama-
zon
Bedrock
Prompt
Man-
age-
ment

Question
23

A logistics company wants to automatically convert audio recordings of customer service calls into written transcripts so that compliance staff can review them as text. Which AWS service is purpose-built for this task?

A.
Ama-
zon
Polly
B.
Ama-
zon
Tran-
scribe
C.
Ama-
zon
Com-
pre-
hend
D.
Ama-
zon
Trans-
late

Question
24

A development team wants their AI agent to be able to securely connect to multiple external tools and data sources, such as a company calendar and a ticketing system, using a standardized way of describing available tools and ex-
chang-

A.
Model
Con-
text
Pro-
tocol
(MCP)

B.
Token-
based
pric-
ing

C.
Diffu-
sion
mod-
eling

D.
Rein-
force-
ment
learn-
ing
from
hu-
man
feed-
back
(RLHF)

Question
25

A
secu-
rity
team
is re-
view-
ing
risks
asso-
ci-
ated
with
a
customer-
facing
GenAI
chat-
bot
built
using
prompt
engi-
neer-
ing.
Se-
lect
TWO
risks
they
should
specif-
ically
eval-
uate
for.

A.
Prompt
injec-
tion/poisoning,
where
mali-
cious
in-
struc-
tions
are
smug-
gled
into
the
model
via
un-
trusted
input
data

B.
Ex-
ces-
sive
use
of
the
tem-
pera-
ture
pa-
rame-
ter
set
to ex-
actly
0.5 C.

Jail-
break-
ing,
where
users
at-
tempt
to
by-
pass
the
model's
safety
re-

*Select
TWO.*

Question

26

A
com-
pany's
legal
team
is as-
sess-
ing
risks
be-
fore
launch-
ing a
GenAI
content-
generation
tool
for
pub-
lic
use.
Se-
lect
THREE
legal
risks
they
should
eval-
uate.

A.
Intel-
lec-
tual
prop-
erty
in-
fringe-
ment
claims
aris-
ing
from
train-
ing
data
or
gen-
er-
ated
out-
puts
B.
Loss
of
cus-
tomer
trust
re-
sult-
ing
from
AI
fail-
ures
or
em-
bar-
rass-
ing
out-
puts
C.
The
num-
ber
of
GPU_s
used
to
train
the

*Select
THREE.*

Question

27

A
com-
pany
is se-
lect-
ing a
foun-
da-
tion
model
for a
new
customer-
facing
appli-
ca-
tion.
Se-
lect
THREE
fac-
tors
that
should
rea-
son-
ably
influ-
ence
the
model
selec-
tion
deci-
sion.

A.
The
model's
la-
tency
(how
quickly
it re-
turns
a re-
sponse)

B.
The
model's
com-
pli-
ance
with
rele-
vant
data-
handling
regu-
la-
tions

C.
The
fa-
vorite
pro-
gram-
ming
lan-
guage
of
the
lead
de-
vel-
oper

D.
The
cost
per
to-
ken
or
cost
per
infer-
ence

E.
The

*Select
THREE.*

Question
28

A
com-
pany
wants
an
AI
sys-
tem
that
can
au-
tonomously
check
in-
ven-
tory
lev-
els,
place
a re-
order
with
a
sup-
plier's
sys-
tem,
and
send
a
con-
fir-
ma-
tion
email,
all
with-
out a
hu-
man
man-
ually
per-
form-
ing
each
step.
Which
AI
63
capa-
bility
best

A. A
sim-
ple
clas-
sifica-
tion
model

B.
An
AI
agent
capa-
ble of
multi-
step
au-
tonomous
task
com-
ple-
tion

C. A
static
rules-
based
spread-
sheet
macro

D. A
single-
shot
prompt
tem-
plate

Question
29

A
sub-
scrip-
tion
stream-
ing
com-
pany
has
years
of
cus-
tomer
data
with
no
pre-
de-
fined
cate-
gories
and
wants
to
dis-
cover
natu-
ral
group-
ings
of
view-
ers
with
simi-
lar
watch-
ing
habits,
with-
out
know-
ing
in ad-
vance
how
many
groups
exist
or
what
de-
c-

-
- A.
Re-
gres-
sion
 - B.
Clas-
sifica-
tion
 - C.
Clus-
ter-
ing
 - D.
Rein-
force-
ment
learn-
ing
-

Question
30

A health-care AI team needs to use real patient records for model development while minimizing the risk of exposing identifiable individual information during that process. Which category of technique addresses this need?

A.
Privacy-
enhancing
tech-
nolo-
gies,
such
as
data
mask-
ing
or
dif-
feren-
tial
pri-
vacy

B. In-
creas-
ing
the
model's
con-
text
win-
dow

C.
Switch-
ing
to a
multi-
modal
model

D.
Us-
ing a
higher
tem-
pera-
ture
set-
ting

Question
31

A
com-
pany's
AI
agent
needs
to re-
mem-
ber a
cus-
tomer's
pref-
er-
ences
across
mul-
tiple
sepa-
rate
con-
versa-
tions
oc-
cur-
ring
days
apart,
rather
than
treat-
ing
each
con-
versa-
tion
as if
it
had
never
hap-
pened
be-
fore.
Which
agen-
tic
AI
con-
cept
ad-
dresses
this

A.
Work-
flow
or-
ches-
tra-
tion
B.
Tool
us-
age
C.
Mem-
ory
man-
age-
ment
D.
Tok-
eniza-
tion

Question
32

A
team
wants
the
model
to
clas-
sify
cus-
tomer
feed-
back
into
“posi-
tive,”
“neu-
tral,”
or
“neg-
a-
tive”
and
pro-
vides
the
model
with
three
ex-
am-
ple
feed-
back
mes-
sages,
each
al-
ready
la-
beled
with
its
cor-
rect
cate-
gory,
be-
fore
ask-
ing it
to
clas-
sify

A.
Zero-
shot
prompt-
ing

B.
Few-
shot
prompt-
ing

C.
Chain-
of-
thought
prompt-
ing

D.
Nega-
tive
prompt-
ing

Question

33

A
team
build-
ing a
facial
recog-
ni-
tion
train-
ing
dataset
wants
to en-
sure
peo-
ple of
many
dif-
fer-
ent
ages,
eth-
nici-
ties,
and
light-
ing
con-
di-
tions
are
well
rep-
re-
sented,
rather
than
skewed
to-
ward
one
nar-
row
group.
Which
re-
spon-
sible
AI
dataset
char-
acter-

A.
Tok-
eniza-
tion
B.
Inclu-
sivity
and
bal-
anced
rep-
re-
sen-
ta-
tion
C.
En-
cryp-
tion
at
rest
D.
Model
distil-
la-
tion

Question
34

A
finan-
cial
ser-
vices
firm
is de-
cid-
ing
be-
tween
a
tradi-
tional
ML
model
and
a
foun-
da-
tion
model
for a
loan-
approval
deci-
sion
sys-
tem
that
must
be
ex-
plain-
able
to
regu-
la-
tors.
Se-
lect
TWO
fac-
tors
that
would
make
a
tradi-
tional
ML
model

A.
The
deci-
sion
must
be
easily
ex-
plain-
able
to a
regu-
lator

B.
The
appli-
ca-
tion
needs
to
gen-
erate
cre-
ative
mar-
ket-
ing
copy

C.
The
orga-
niza-
tion
has
strict
re-
quire-
ments
to
trace
ex-
actly
how
each
input
fea-
ture
af-
fected
the
out-
come

D.

*Select
TWO.*

Question
35

A
com-
pany's
cus-
tomer
sup-
port
team
is im-
pressed
that
their
new
GenAI
assis-
tant
can
han-
dle
ques-
tions
about
re-
turns,
ship-
ping,
prod-
uct
specs,
and
ac-
count
is-
sues,
all
using
natu-
ral
con-
versa-
tional
lan-
guage,
with-
out
need-
ing a
sepa-
rate
78
pro-
gram
built
f

A.
De-
ter-
min-
ism
B.
Adapt-
abil-
ity
across
a
wide
range
of
tasks
and
do-
mains
using
natu-
ral
lan-
guage
C.
Guar-
an-
teed
fac-
tual
accu-
racy
D.
Zero
oper-
ating
cost

Question
36

A
team's
prompts
are
pro-
duc-
ing
in-
con-
sis-
tent,
overly
long,
and
some-
times
irrel-
evant
re-
sponses
from
their
foun-
da-
tion
model.
Which
prompt
engi-
neer-
ing
best
prac-
tice
would
most
di-
rectly
help
ad-
dress
this
is-
sue?

A.
Make
the
prompt
as
vague
as
possi-
ble
to
give
the
model
maxi-
mum
free-
dom
B.
Write
a
clear,
spe-
cific,
and
ap-
pro-
pri-
ately
con-
cise
prompt
with
well-
defined
in-
struc-
tions
C.
Re-
move
all
con-
text
from
the
prompt
to re-
duce
to-
ken
us-
age
D.

Question
37

A
com-
pany
wants
its
GenAI
appli-
ca-
tion
to
an-
chor
every
fac-
tual
claim
it
makes
to a
veri-
fied
source
docu-
ment,
re-
duc-
ing
the
chance
that
the
model
in-
vents
infor-
ma-
tion
not
actu-
ally
sup-
ported
by
com-
pany
data.
Which
tech-
nique
83
most
di-
rectly

A.
In-
creas-
ing
the
tem-
pera-
ture
pa-
rame-
ter B.
RAG
ground-
ing
C.
Re-
duc-
ing
the
num-
ber
of
fine-
tuning
ex-
am-
ples
D.
Dis-
abling
out-
put
filter-
ing

Question
38

A
com-
pany
de-
ployed
a
GenAI
writ-
ing
assis-
tant
for
its
mar-
ket-
ing
team
six
months
ago.
Lead-
er-
ship
wants
to
know
whether
the
in-
vest-
ment
is
pay-
ing
off in
busi-
ness
terms.
Which
met-
ric
would
most
di-
rectly
help
an-
swer
that
ques-
tion?

A.
The
num-
ber
of pa-
rame-
ters
in
the
un-
derly-
ing
foun-
da-
tion
model

B.
Re-
turn
on in-
vest-
ment
(ROI),
com-
par-
ing
cost
sav-
ings
and
rev-
enue
im-
pact
against
the
tool's
cost

C.
The
model's
con-
text
win-
dow
size
in to-
kens

D.
The
diffu-
sion

Question

39

A
com-
pany
is
prepar-
ing a
dataset
to
fine-
tune
a
foun-
da-
tion
model
for
inter-
nal
cus-
tomer
ser-
vice
use.
Se-
lect

THREE

prac-
tices
that
rep-
re-
sent
good
data
prepa-
ra-
tion
for
fine-
tuning.

A. Cu-
rat-
ing
high-
quality,
rele-
vant
ex-
am-
ples
rather
than
in-
clud-
ing
every
avail-
able
record
re-
gard-
less
of
qual-
ity B.
En-
sur-
ing
the
train-
ing
data
is
prop-
erly
li-
censed
and
legally
us-
able
C. In-
clud-
ing
only
data
from
a sin-
gle⁸⁸
cus-
tomer

*Select
THREE.*

Question
40

A
game
de-
vel-
op-
ment
stu-
dio
wants
an
AI
agent
that
learns
to
play
a
new
game
by re-
peat-
edly
at-
tempt-
ing
moves,
re-
ceiv-
ing a
posi-
tive
score
when
it
suc-
ceeds
and a
penalty
when
it
fails,
grad-
ually
im-
prov-
ing
its
strat-
egy
over
time.
Which

-
- A.
Su-
per-
vised
learn-
ing
 - B.
Un-
su-
per-
vised
learn-
ing
 - C.
Rein-
force-
ment
learn-
ing
 - D.
Trans-
fer
learn-
ing
-

Question
41

A
model
per-
forms
ex-
tremely
well
on
its
train-
ing
data
but
per-
forms
poorly
on
new,
un-
seen
data
in
pro-
duc-
tion.
What
does
this
most
likely
indi-
cate?

-
- A.
Under-
fit-
ting
 - B.
Over-
fit-
ting
 - C.
High
bias
with
no
vari-
ance
 - D.
Per-
fect
gen-
eral-
iza-
tion
-

Question
42

A development team is deciding which AWS offerings to use for building and then operating an AI agent in production. Select TWO statements that correctly describe the distinction between Strands Agents and Amazon Bedrock Agent-Core.

A.
Strands
Agents
is an
open-
source
frame-
work
used
to
build
agents;
Agent-
Core
is the
man-
aged
ser-
vice
used
to de-
ploy,
se-
cure,
and
scale
agents
in
pro-
duc-
tion
B.
Agent-
Core
and
Strands
Agents
are
sim-
ply
two
dif-
fer-
ent
names
for
the
exact
same
95
offer-
ing C.
Strands
Agents

*Select
TWO.*

Question
43

A
com-
pany
has
mul-
tiple
teams
edit-
ing
and
ex-
peri-
ment-
ing
with
prompts
for
the
same
pro-
duc-
tion
GenAI
appli-
ca-
tion,
and
lead-
er-
ship
is
con-
cerned
about
los-
ing
track
of
which
prompt
ver-
sion
is
cur-
rently
live
ver-
sus
which
are
in
test-
.

A.
Ama-
zon
Sage-
Maker
Clar-
ify B.
Ama-
zon
Bedrock
Prompt
Man-
age-
ment
C.
AWS
Au-
dit
Man-
ager
D.
Ama-
zon
Ma-
cie

Question
44

A security team is building a checklist of AI-specific security and privacy considerations to address before launching a customer-facing GenAI application. Select THREE items that belong on this AI-specific checklist.

A.
Prompt
injec-
tion
de-
fenses,
to
pre-
vent
mali-
cious
in-
struc-
tions
hid-
den
in
user
input
or
exter-
nal
data

B.
Out-
put
filter-
ing
and
vali-
da-
tion,
to
catch
inap-
pro-
pri-
ate
or
incor-
rect
con-
tent
be-
fore
it
reaches
users

C.
Toxi-
city
de-

*Select
THREE.*

Question
45

A
hos-
pital
is de-
ploy-
ing
an
AI
model
to
flag
po-
ten-
tially
can-
cer-
ous
scans
for
fur-
ther
re-
view
by a
radi-
olo-
gist.
Miss-
ing
an
ac-
tual
can-
cer
case
is far
more
costly
than
send-
ing a
healthy
scan
for
un-
nec-
es-
sary
extra
re-
view.
Which

A.
Pre-
ci-
sion
B.
Re-
call
C. F1
score
only,
ig-
nor-
ing
preci-
sion
and
recall
indi-
vidu-
ally
D.
Accu-
racy

Question
46

A
small
busi-
ness
with
no in-
house
ML
engi-
neers
wants
to
add a
GenAI-
powered
chat-
bot
to its
web-
site
quickly
and
af-
ford-
ably.
What
is a
key
ad-
van-
tage
of
using
a
man-
aged
AWS
GenAI
ser-
vice
like
Ama-
zon
Bedrock
over
build-
ing a
foun-
da-
tion
model
from

A.
Build-
ing
from
scratch
is al-
ways
cheaper
for
small
busi-
nesses

B.
Man-
aged
ser-
vices
pro-
vide
lower
barri-
ers
to en-
try,
faster
time
to
mar-
ket,
and
re-
duced
oper-
a-
tional
over-
head
com-
pared
to
train-
ing a
model
from
scratch

C.
Man-
aged
ser-
vices
re-
quire

Question
47

A
foun-
da-
tion
model
provider
peri-
odi-
cally
up-
dates
its
model
by
feed-
ing it
new,
broad
data
over
time
to
keep
its
gen-
eral
knowl-
edge
cur-
rent,
with-
out
tar-
get-
ing
any
sin-
gle
spe-
cific
task
or
do-
main.
Which
train-
ing
stage
does
this
de-
scribe?

A.
Fine-
tuning

B.
Con-
tinu-
ous
pre-
training

C. In-
struc-
tion
tun-
ing

D.
Model
distil-
la-
tion

Question
48

An organization wants to both detect bias in a trained model and incorporate human review for predictions where the model has low confidence. Select TWO AWS tools that directly support these two distinct needs, respectively.

A.
Ama-
zon
Sage-
Maker
Clar-
ify,
for
de-
tect-
ing
bias
in
mod-
els
and
datasets

B.
Ama-
zon
Aug-
mented
AI
(A2I),
for
rout-
ing
low-
confidence
pre-
dic-
tions
to
hu-
man
re-
view-
ers C.

AWS
Cloud-
Trail,
for
de-
tect-
ing
bias
in
datasets

D.
110
Ama-
zon
Ma-
.

*Select
TWO.*

Question
49

A health-care company is choosing between building its own GenAI infrastructure from scratch or using AWS-managed GenAI services, partly because it needs to meet strict health-care compliance requirements. Which benefit of AWS in-
112
fras-
truc-
ture

A.
AWS
in-
fras-
truc-
ture
in-
cludes
built-
in
secu-
rity
certi-
fica-
tions,
com-
pli-
ance
capa-
bili-
ties,
and
clearly
de-
fined
re-
spon-
sibil-
ity
bound-
aries
that
sup-
port
regu-
lated
work-
loads

B.
AWS
in-
fras-
truc-
ture
guar-
an-
tees
the
com-
pany
will
never

Question
50

A
team
is de-
sign-
ing
an
eval-
ua-
tion
strat-
egy
for a
new
customer-
facing
foun-
da-
tion
model
appli-
ca-
tion
be-
fore
launch.
Se-
lect
TWO
eval-
ua-
tion
ap-
proaches
that
are
docu-
mented
as
ap-
pro-
pri-
ate
meth-
ods
for
as-
sess-
ing
FM
per-
for-
mance.

A.
Human-
in-
the-
loop
eval-
ua-
tion,
where
peo-
ple
rate
the
qual-
ity of
model
out-
puts

B.
Com-
par-
ing
out-
puts
against
bench-
mark
datasets
with
known
refer-
ence
an-
swers

C.
Ask-
ing
the
origi-
nal
model
itself,
with
no
other
vali-
da-
tion,
whether
its
own
an-

*Select
TWO.*

Question
51

A company wants to download AWS's own existing compliance certifications and audit reports (such as SOC 2 reports) to share with its own auditors, without having to generate any new evidence itself. Which AWS service should they use?

A.
AWS
Au-
dit
Man-
ager
B.
AWS
Arti-
fact
C.
AWS
Con-
fig D.
Ama-
zon
In-
spec-
tor

Question
52

A
com-
pany
is
map-
ping
its
planned
AI
ini-
tia-
tives
to AI
appli-
ca-
tion
cate-
gories.
Se-
lect
the
THREE
ini-
tia-
tives
that
are
cor-
rectly
matched
to
their
cate-
gory.

A.
Us-
ing
AI to
flag
un-
usual
credit
card
trans-
ac-
tions
in
real
time

-
fraud
de-
tec-
tion

B.
Us-
ing
AI to
rec-
om-
mend
prod-
ucts
based
on
brows-
ing
his-
tory -
rec-
om-
men-
da-
tion
sys-
tem

C.
Us-
ing
AI to
pre-
dict
next
121
quarter's
re-

*Select
THREE.*

Question
53

A
com-
pany
has a
GenAI
appli-
ca-
tion
with
highly
pre-
dictable,
con-
sis-
tently
high
re-
quest
vol-
ume
through-
out
the
day
and
wants
guar-
an-
teed,
con-
sis-
tent
re-
sponse
times
rather
than
vari-
able
per-
for-
mance
dur-
ing
peak
de-
mand.
Which
Ama-
zon
Bedrock
pric-

-
- A.
On-
demand
token-
based
pric-
ing
only
- B.
Pro-
vi-
sioned
through-
put
- C.
Spot
pric-
ing
for
foun-
da-
tion
mod-
els
- D.
Free-
tier
ac-
cess
only
-

Question
54

A
com-
pany
is
eval-
uat-
ing
an
FM-
powered
tool
that
pro-
duces
trans-
lated
prod-
uct
de-
scrip-
tions
from
En-
glish
into
French
and
wants
an
auto-
mated
met-
ric
specif-
ically
suited
to as-
sess-
ing
trans-
la-
tion
qual-
ity.
Which
met-
ric is
most
ap-
pro-
pri-
ate?

-
- A.
ROUGE
 - B.
BLEU
 - C. F1
score
 - D.
Accu-
racy
-

Question
55

A
com-
pany
wants
to
pro-
duce
struc-
tured
docu-
men-
ta-
tion
de-
scrib-
ing a
model's
in-
tended
use,
per-
for-
mance
char-
acter-
istics,
and
known
limi-
ta-
tions,
to
share
with
stake-
hold-
ers
and
audi-
tors.
Which
AWS
tool
is de-
signed
for
this
pur-
pose?

A.
Ama-
zon
Sage-
Maker
Model
Cards
B.
AWS
Trusted
Advi-
sor C.
Ama-
zon
Polly
D.
AWS
Glue

Question
56

A
data
scien-
tist
wants
to
quickly
de-
ploy
a pre-
built
foun-
da-
tion
model
di-
rectly
from
within
the
Sage-
Maker
inter-
face,
as
part
of a
broader
cus-
tom
ML
work-
flow
she's
al-
ready
build-
ing
in
Sage-
Maker.
Which
ser-
vice
is
most
aligned
with
this
need?

A.
Ama-
zon
Bedrock
Agent-
Core
B.
Ama-
zon
Sage-
Maker
Jump-
Start
C.
AWS
Au-
dit
Man-
ager
D.
Ama-
zon
Ma-
cie

Question
57

A
team
has
thou-
sands
of
chat-
bot
re-
sponses
to
eval-
uate
for
qual-
ity
and
tone
but
does
not
have
the
staffing
to
man-
ually
re-
view
each
one.
They
want
a
scal-
able
auto-
mated
method
that
goes
be-
yond
sim-
ple
word-
overlap
met-
rics.
Which
eval-
ua-
ti-

A.
ROUGE
score
only

B.
LLM-
as-a-
judge,
using
a sep-
arate
trusted
model
to
eval-
uate
the
qual-
ity of
the
out-
puts

C.
Man-
ual
re-
view
by
the
origi-
nal
engi-
neer-
ing
team
of
every
sin-
gle
re-
sponse

D.
BLEU
score
only

Question
58

A retail demand-forecasting model was highly accurate when launched, but six months later its predictions have become noticeably less reliable as customer buying patterns shifted post-launch. According to MLOps best practices, what should the team do?

A.
De-
com-
mis-
sion
the
model
per-
ma-
nently,
since
AI
mod-
els
can-
not
adapt
to
chang-
ing
con-
di-
tions

B.
Con-
tinue
using
the
model
un-
changed,
since
re-
train-
ing is
only
needed
if the
origi-
nal
code
has
bugs

C.
Mon-
itor
the
model's
per-
for-
mance
and

Question
59

A multi-national company must ensure that AI training data for its European customers remains stored only within European AWS regions, due to local regulatory requirements. Which data governance concept does this require-
137
ment re-

-
- A.
Data
lin-
eage
 - B.
Data
resi-
dency
 - C.
Data
cata-
loging
 - D.
Token-
based
pric-
ing
-

Question
60

An AI practitioner is evaluating two candidate models for a customer support chatbot. Model A is 3% more accurate than Model B but costs four times as much per month to run, and the support team reports no noticeable

A.
Model
A
should
auto-
mati-
cally
be
cho-
sen
be-
cause
higher
accu-
racy
al-
ways
means
higher
busi-
ness
value

B.
Model
B
may
be
the
bet-
ter
busi-
ness
choice,
since
the
marginal
accu-
racy
gain
from
Model

A
does
not
ap-
pear
to
trans-
late
into
bet-
ter
business

Question

61

A
com-
pany
has
de-
ployed
an
AI
agent
that
han-
dles
cus-
tomer
sup-
port
tick-
ets
end-
to-
end.
Se-
lect
TWO
met-
rics
that
would
help
de-
ter-
mine
whether
the
agent
is
meet-
ing
busi-
ness
ob-
jec-
tives.

A.
Task
com-
ple-
tion
rate -
the
per-
cent-
age
of
sup-
port
tick-
ets
the
agent
suc-
cess-
fully
re-
solves
with-
out
esca-
la-
tion

B.
Cost
per
inter-
ac-
tion -
the
total
cost
to
han-
dle
one
sup-
port
ticket
using
the
agent

C.
The
num-
ber
of
lines

*Select
TWO.*

Question
62

A
com-
pany
de-
ploys
an
AI
tool
that
auto-
mati-
cally
flags
cer-
tain
loan
appli-
ca-
tions
for
addi-
tional
re-
view.
To
sup-
port
human-
centered
de-
sign
prin-
ci-
ples,
what
should
the
com-
pany
also
pro-
vide
to
appli-
cants
whose
loans
were
flagged?

A.
No
ex-
pla-
na-
tion
at all,
since
the
model's
inter-
nal
logic
is
pro-
pri-
etary

B. A
way
for
appli-
cants
to
un-
der-
stand
that
AI
played
a
role
in
the
deci-
sion
and a
mech-
a-
nism
to
pro-
vide
feed-
back
or
raise
con-
cerns

C.⁴⁵
The
com-
plete

Question

63

A
com-
pany
wants
to
add
a fea-
ture
to its
mo-
bile
app
that
reads
order
con-
fir-
ma-
tion
text
aloud
to
visu-
ally
im-
paired
cus-
tomers
in a
natural-
sounding
voice.
Which
AWS
ser-
vice
should
they
use?

A.
Ama-
zon
Lex
B.
Ama-
zon
Polly
C.
Ama-
zon
Tran-
scribe
D.
Ama-
zon
Com-
pre-
hend

Question
64

A
com-
pany
wants
to
use a
rec-
og-
nized
frame-
work
to
help
clas-
sify
its
vari-
ous
GenAI
use
cases
by
the
level
of
risk
and
re-
spon-
sibil-
ity
in-
volved,
in
order
to
apply
ap-
pro-
pri-
ately
scaled
gov-
er-
nance
con-
trols
to
each
one.
Which
gov-

A.
The
AWS
Well-
Architected
Frame-
work

B.
The
Gen-
era-
tive
AI
Secu-
rity
Scop-
ing
Ma-
trix

C.
The
AWS
Shared
Re-
spon-
sibil-
ity
Model

D.
ROUGE
scor-
ing

Question
65

A
com-
pany
has a
very
large,
highly
accu-
rate
foun-
da-
tion
model
but
needs
a
smaller,
faster,
less
ex-
pen-
sive
ver-
sion
to
run
on
resource-
constrained
de-
vices,
while
re-
tain-
ing
as
much
of
the
origi-
nal
model's
be-
hav-
ior as
possi-
ble.
Which
tech-
nique
best
fits
this

A.
Model
distil-
la-
tion
B.
RAG
C.
Prompt
engi-
neer-
ing
D.
Data
label-
ing

End of Section 2. Record your answers on the Answer Sheet below.

SECTION 3: ANSWER SHEET

Q#	Answer	Q#	Answer	Q#	Answer	Q#	Answer	Q#	Answer
1		14		27		40		53	
2		15		28		41		54	
3		16		29		42		55	
4		17		30		43		56	
5		18		31		44		57	
6		19		32		45		58	
7		20		33		46		59	
8		21		34		47		60	
9		22		35		48		61	
10		23		36		49		62	
11		24		37		50		63	
12		25		38		51		64	
13		26		39		52		65	

For multiple-response questions, write all correct letters (for example, “A, C”). For the ordering questions, write your full sequence (for example, “2, 4, 1, 5, 3”). For the matching question, write each pairing (for example, “1-C, 2-B, 3-A”).

SECTION 4: DOMAIN PERFORMANCE TRACKER

After checking your answers against the Answer Key (Practice_Exam_With_Answers.md), record your scores below.

Domain	Questions	# Correct	Score %	Action
Domain 1: Fundamentals of AI and ML	13			
Domain 2: Fundamentals of Generative AI	16			
Domain 3: Applications of Foundation Models	18			
Domain 4: Guidelines for Responsible AI	9			
Domain 5: Security, Compliance, and Governance	9			
TOTAL	65			

Scoring guide: 47+ correct (73%+) = on track to pass | Below 47 = continue studying before scheduling

End of Task 10 - Full-Length Practice Exam, Learner Version. Document based on AIF-C01 Exam Guide Version 1.1 (April 30, 2026).